

AWS  
re:Invent

DAT305

# Real-world use cases for Amazon DynamoDB

**Sean Shriver**

Sr. DynamoDB SA  
Amazon Web Services

**Ankur Kasliwal**

Technical Program Manager  
Amazon Web Services

# Three things to consider



Hi!



URL



Office hours

# Agenda

1. DynamoDB fundamentals
2. Three distilled steps for a successful design
3. Three real-world use cases for DynamoDB

# Amazon DynamoDB



Key value



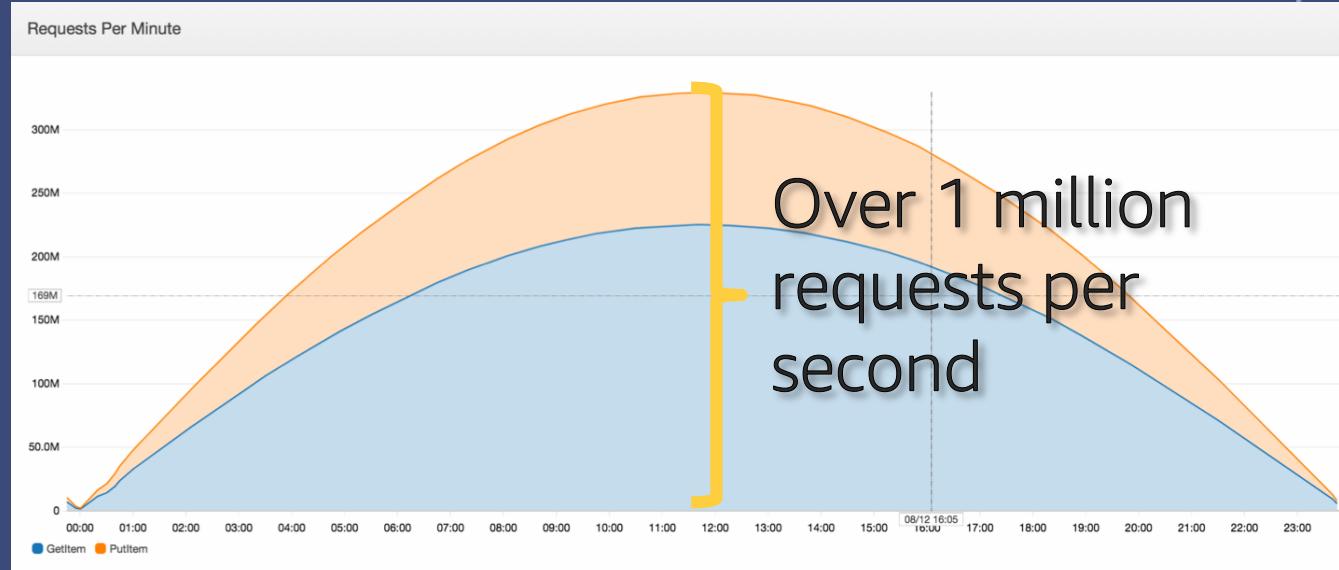
Durable



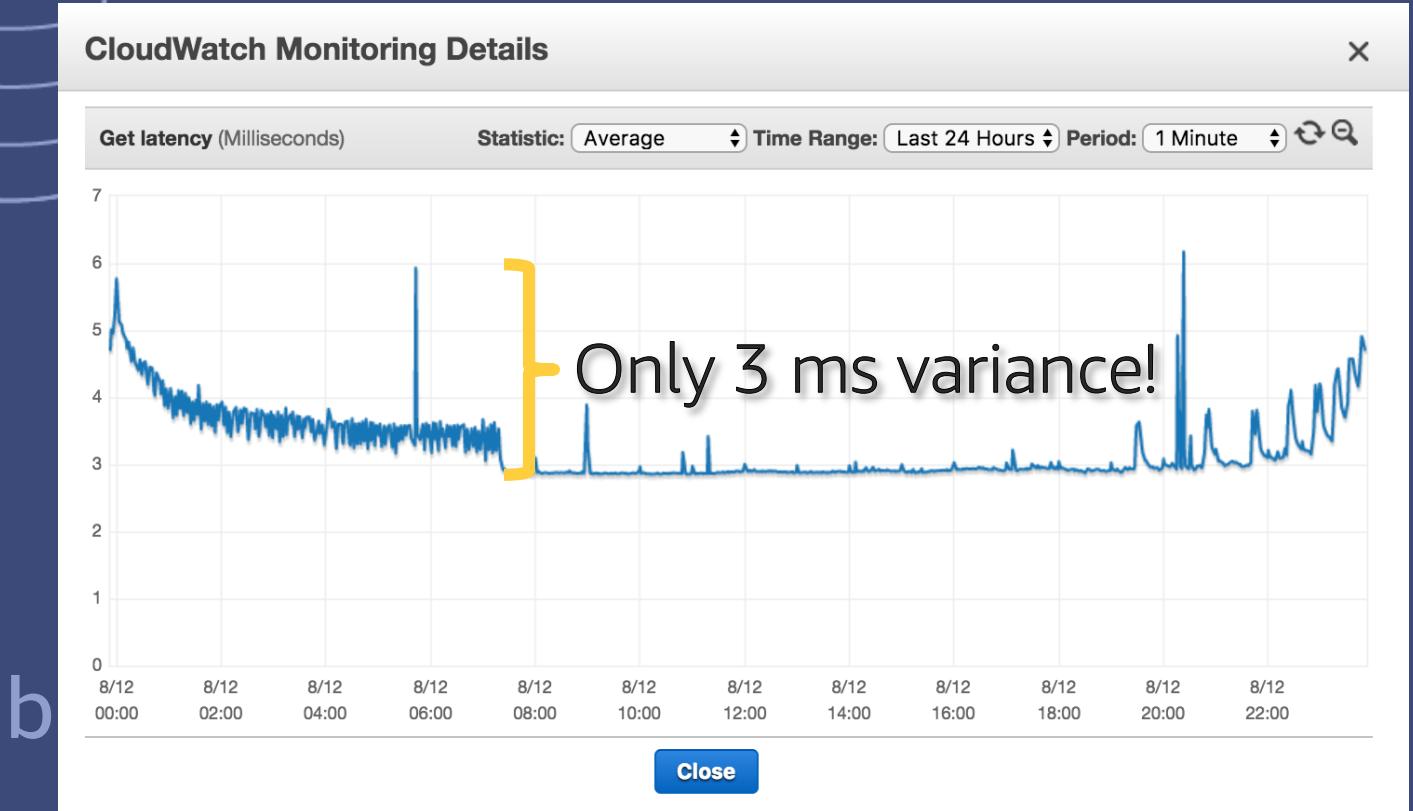
Low latency

# Amazon DynamoDB

## Request Volume



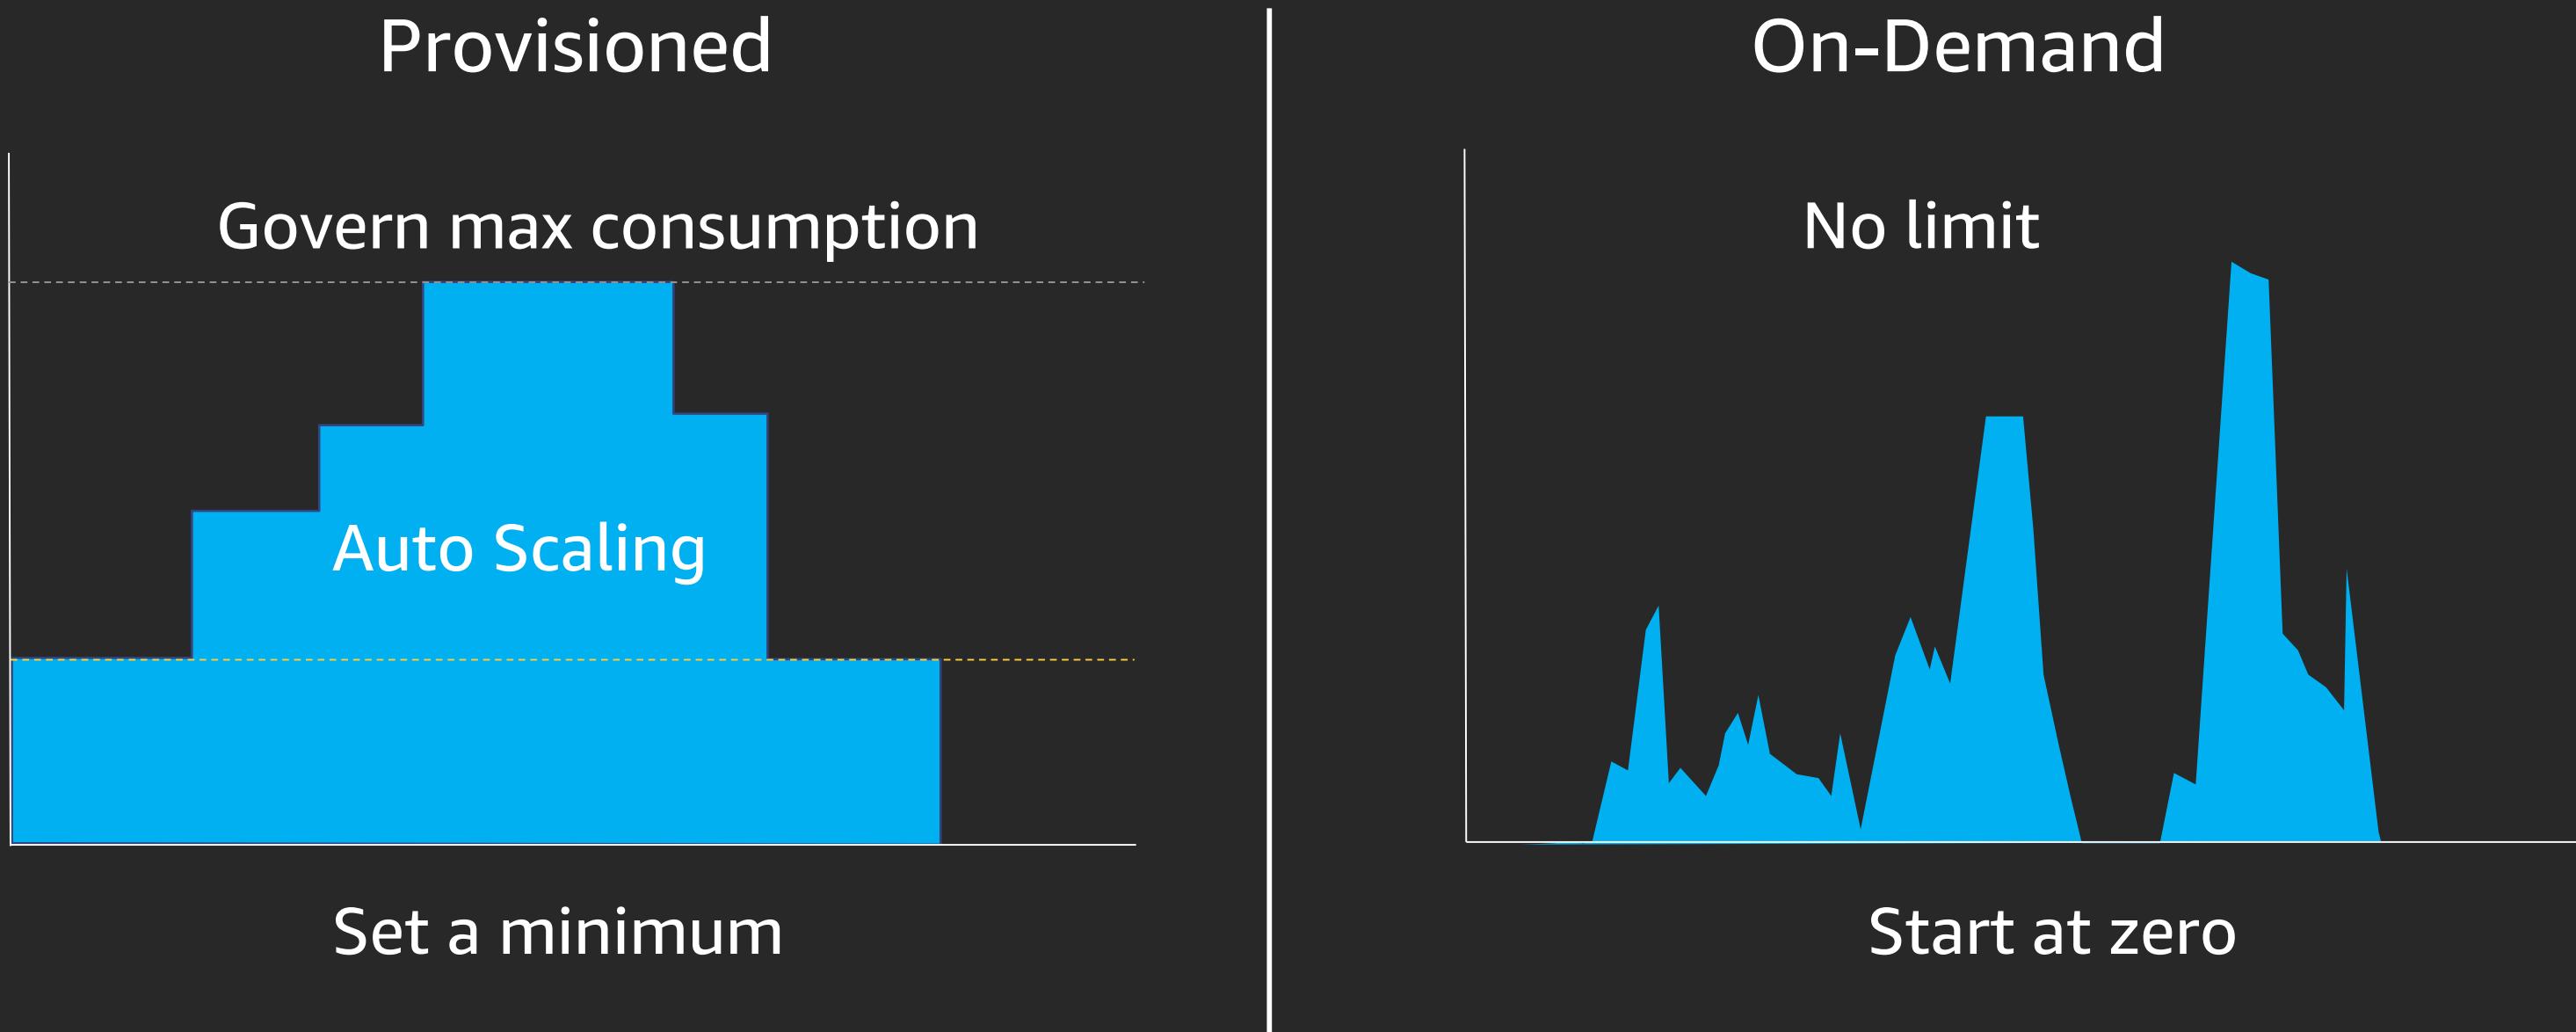
## Latency



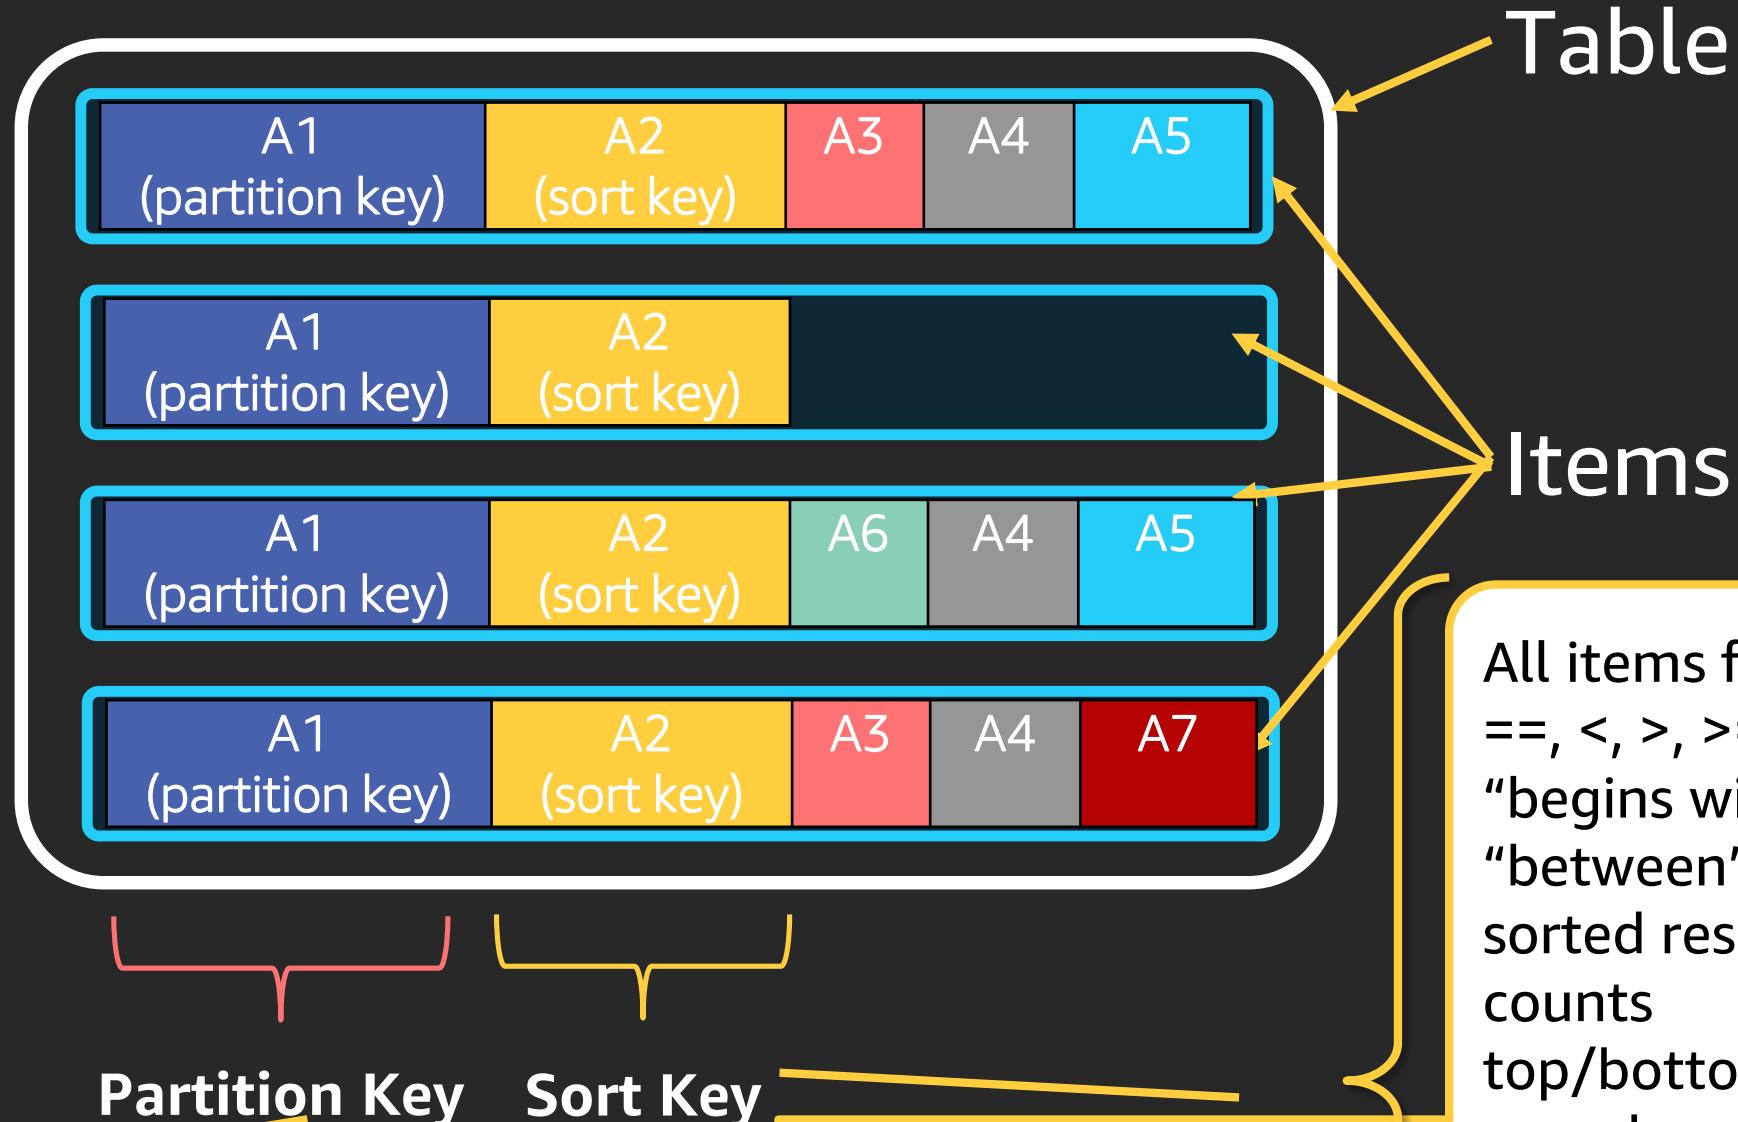
Key Value

Durab

# DynamoDB: Capacity managed for you



# DynamoDB table



Mandatory  
Key-value access pattern  
Determines data distribution

Optional  
Model 1:N relationships  
Enables rich query capabilities

All items for a partition key  
==, <, >, >=, <=  
“begins with”  
“between”  
sorted results  
counts  
top/bottom N values  
paged responses

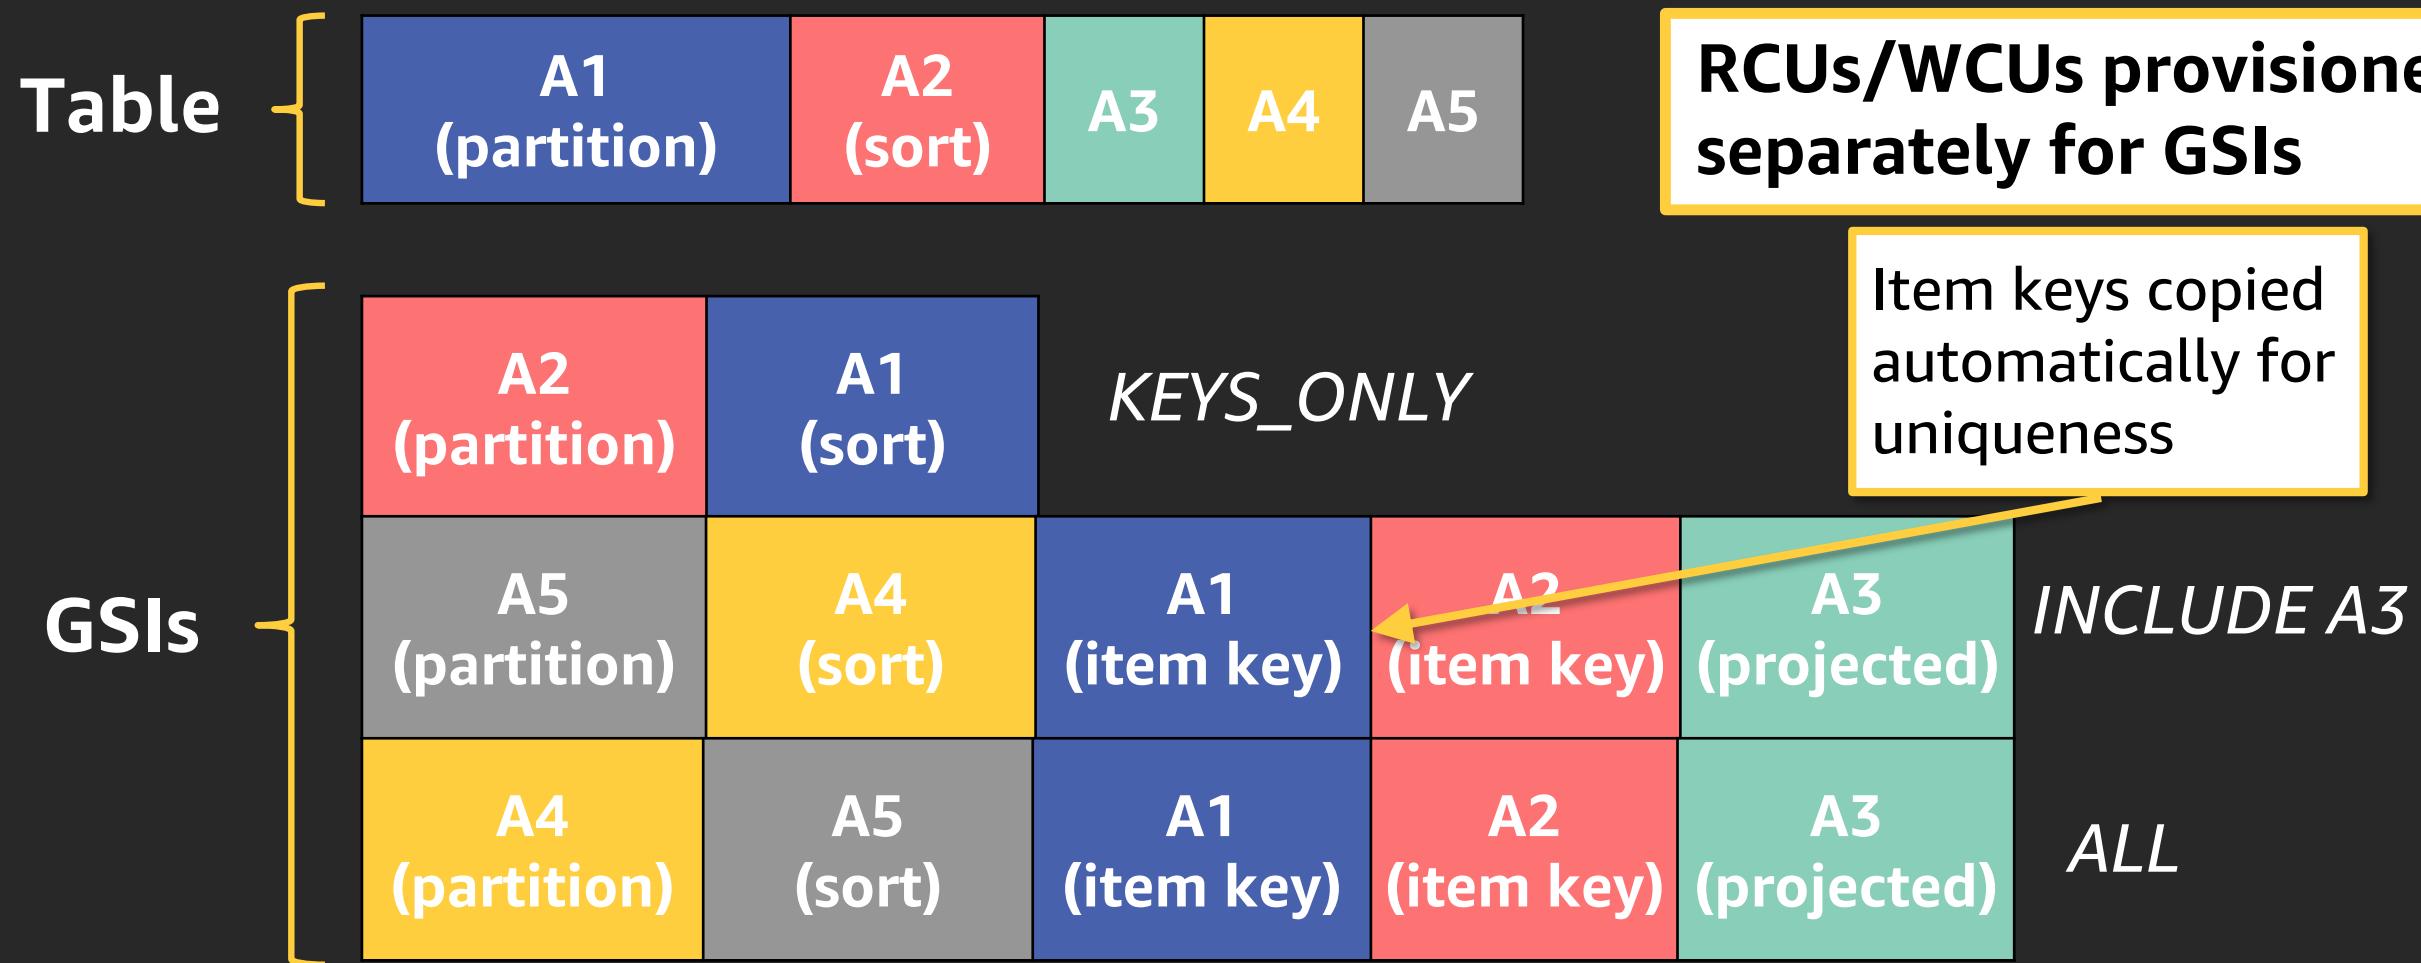
# Global secondary index (GSI)

Up to 20 GSIs per table

Alternate partition and/or sort key

Index is across all partition keys

Use composite sort keys for compound indexes



# Local secondary index (LSI)

Alternate sort key attribute

Index is local to a partition key

Table

A1 (partition)	A2 (sort)	A3	A4	A5
-------------------	--------------	----	----	----

## Collection limits

10 GB max per partition key.  
LSIs limit the # of range keys.

LSIs

A1 (partition)	A3 (sort)	A2 (item key)	KEYS_ONLY	
A1 (partition)	A4 (sort)	A2 (item key)	A3 (projected)	INCLUDE A3
A1 (partition)	A5 (sort)	A2 (item key)	A3 (projected)	A4 (projected)

ALL

# Distilled design concepts

# Distilled design concepts

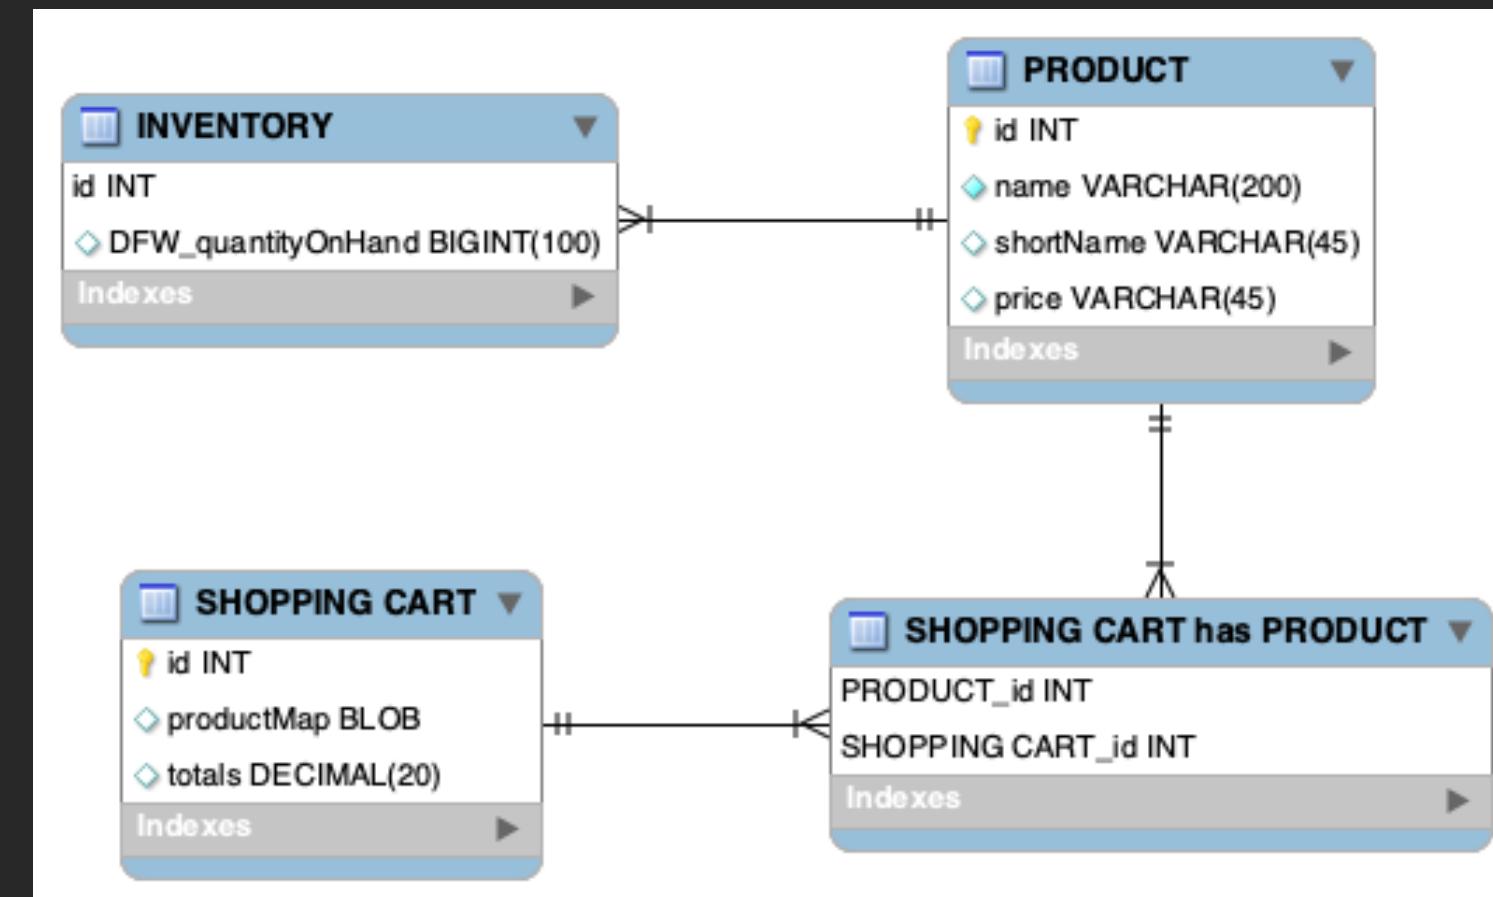
1. Understand the use case
2. Model the entity relationships
3. Write down the queries

# 1. Understand the use case

- ✓ Industry knowledge
- ✓ Is this a new or existing feature?
- ✓ Is DynamoDB the operational database?
- ✓ Who are the actors querying data from DynamoDB?

## 2. Model the relationships

- ✓ Entity-relationship diagrams are here to stay



### 3. Establish the queries

- ✓ Get all PRODUCTS from SHOPPING CART
- ✓ Get all SHOPPING CARTs holding a PRODUCT
- ✓ Reserve the PRODUCT in INVENTORY during checkout

Tip: think in SQL, and then convert to words

### 3. Establish the queries

- ✓ Get all PRODUCTS from SHOPPING CART

```
SELECT products FROM cart where cart_id = 'YYY';
```

- ✓ Get all SHOPPING CARTs holding a PRODUCT

```
SELECT cart_id FROM products_cart WHERE product_id = 'XXX';
```

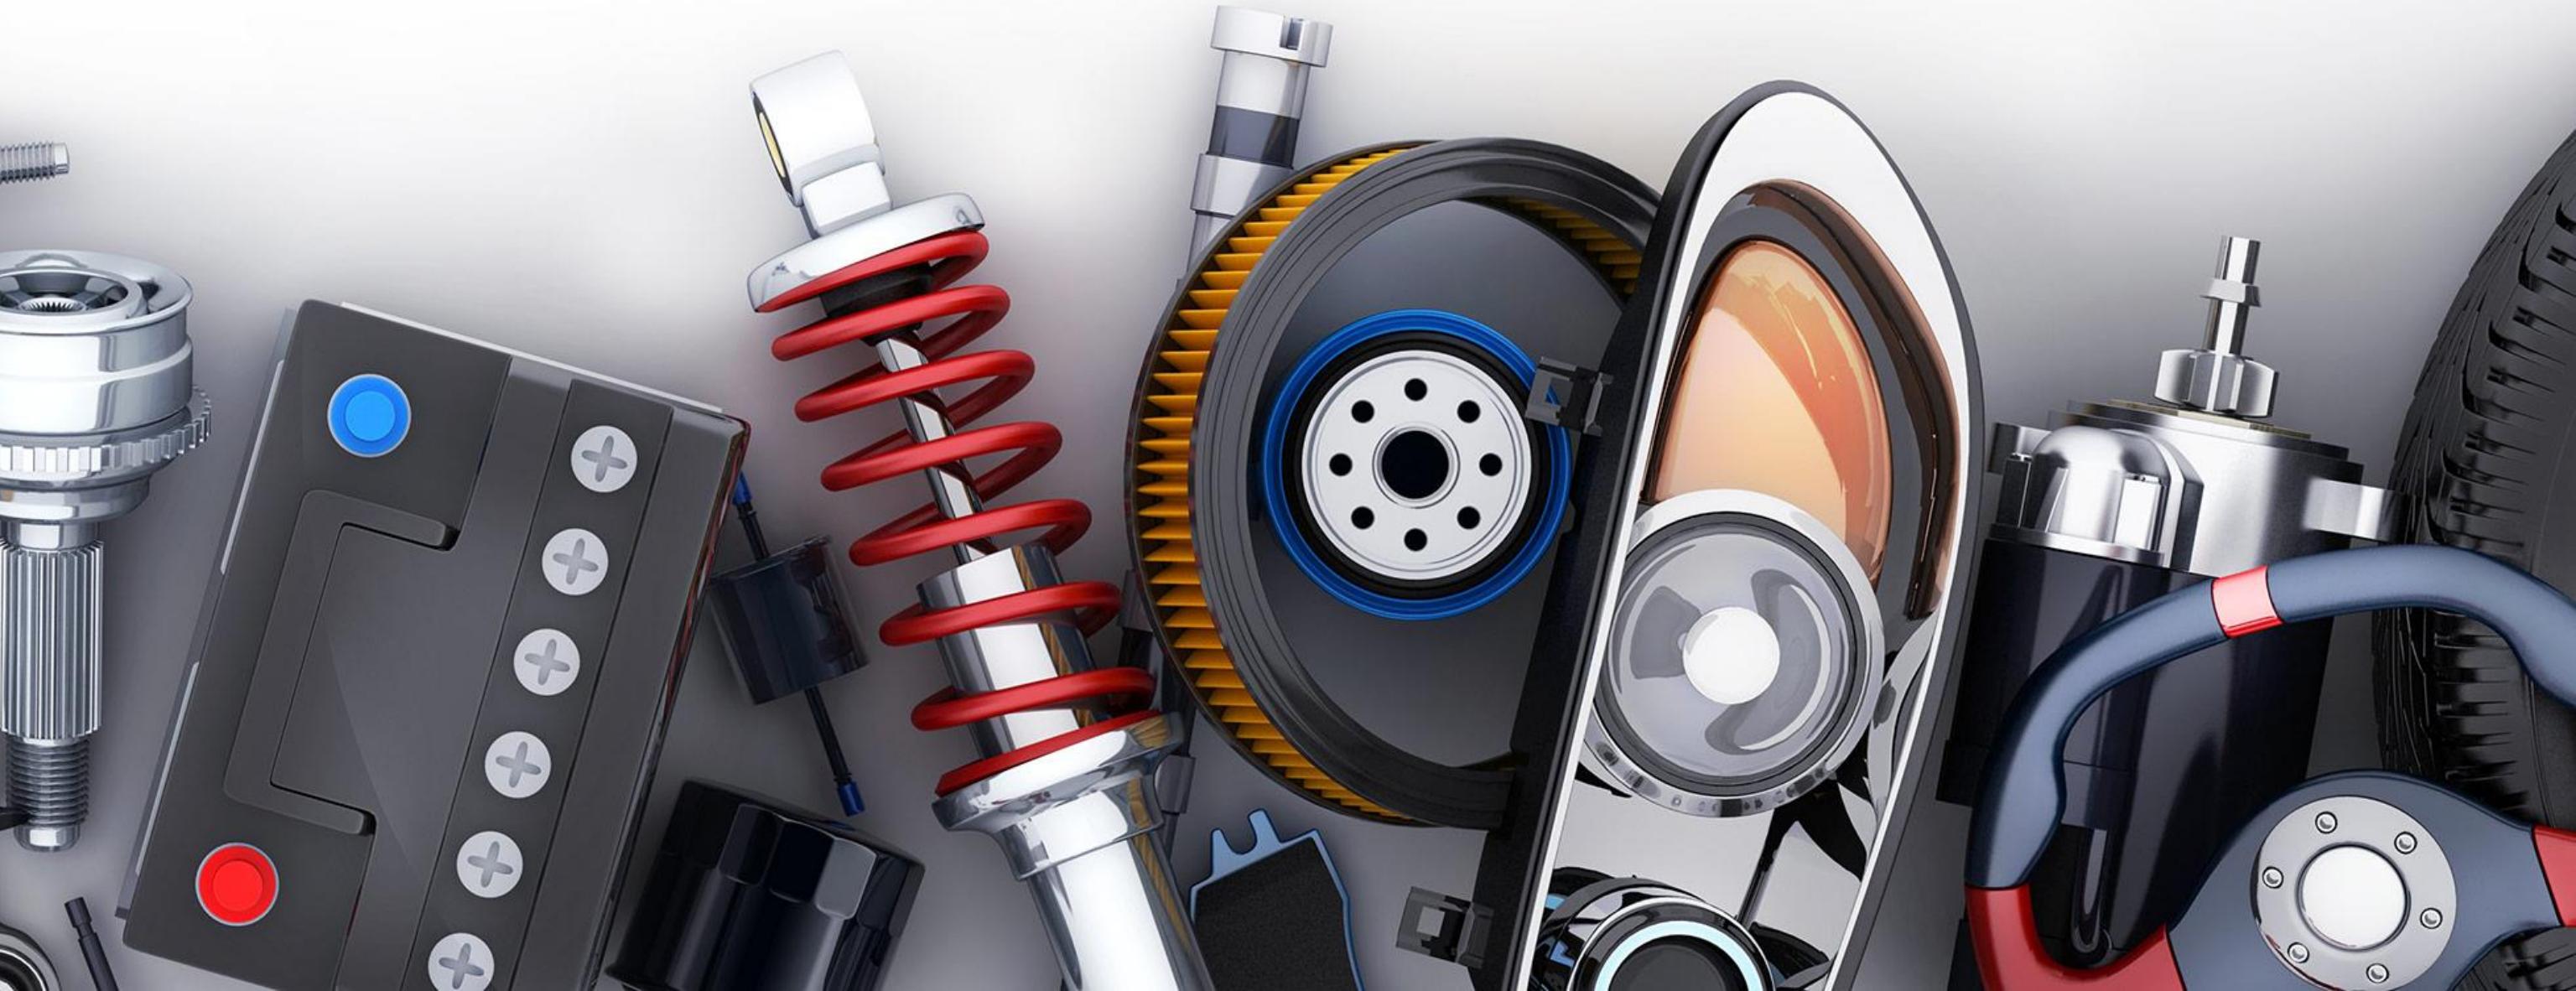
- ✓ Reserve the PRODUCT in INVENTORY during checkout

```
UPDATE inventory SET qoh = qoh - 1 WHERE product_id = 'XXX'  
AND qoh > 0;
```

Remember, DynamoDB is API driven and does not use SQL.

# Real world use cases

Customer #1: Midwest industrial manufacturer  
Use case: *n:m relationships, data lake processing*



Why Dynamo?  
Low latency  
Need durability  
Operational data

Schedule

8A - Start  
10-noon - Lab  
Noon-1P - EAT  
1P-3P - Design



Possible partition keys:  
Process order id  
Possible sort keys:  
???

Queries:

1. Get all steps for  
process order



Why Dynamo?  
Low latency  
Need durability  
Operational data

Schedule

8A - Start  
10-noon - Lab  
Noon-1P - EAT  
1P-3P - Design



Possible partition keys:  
Process order id  
Possible sort keys:  
???

Queries:  
1. Get all steps for process order

## Why Dynamo?

- ✓ Low latency
- ✓ Need durability
- ✓ Operational data

## Schedule

8A - Start  
10-noon - Lab  
Noon-1P - EAT  
1P-3P - Design



## Possible partition keys:

- ✓ Process order id

## Possible sort keys:

???

## Queries:

1. Get all steps for process order

# Real-world additive manufacturing

*Process an order for a set of parts, and track it all!*

Step 1: Gather  
“Materials”

*e.g. Aluminum*

Step 2: Mix  
“Materials”

*Mix, and measure  
mass and volume*

Step 3: Print  
“Parts”

*2 or more parts  
e.g. drill bit*

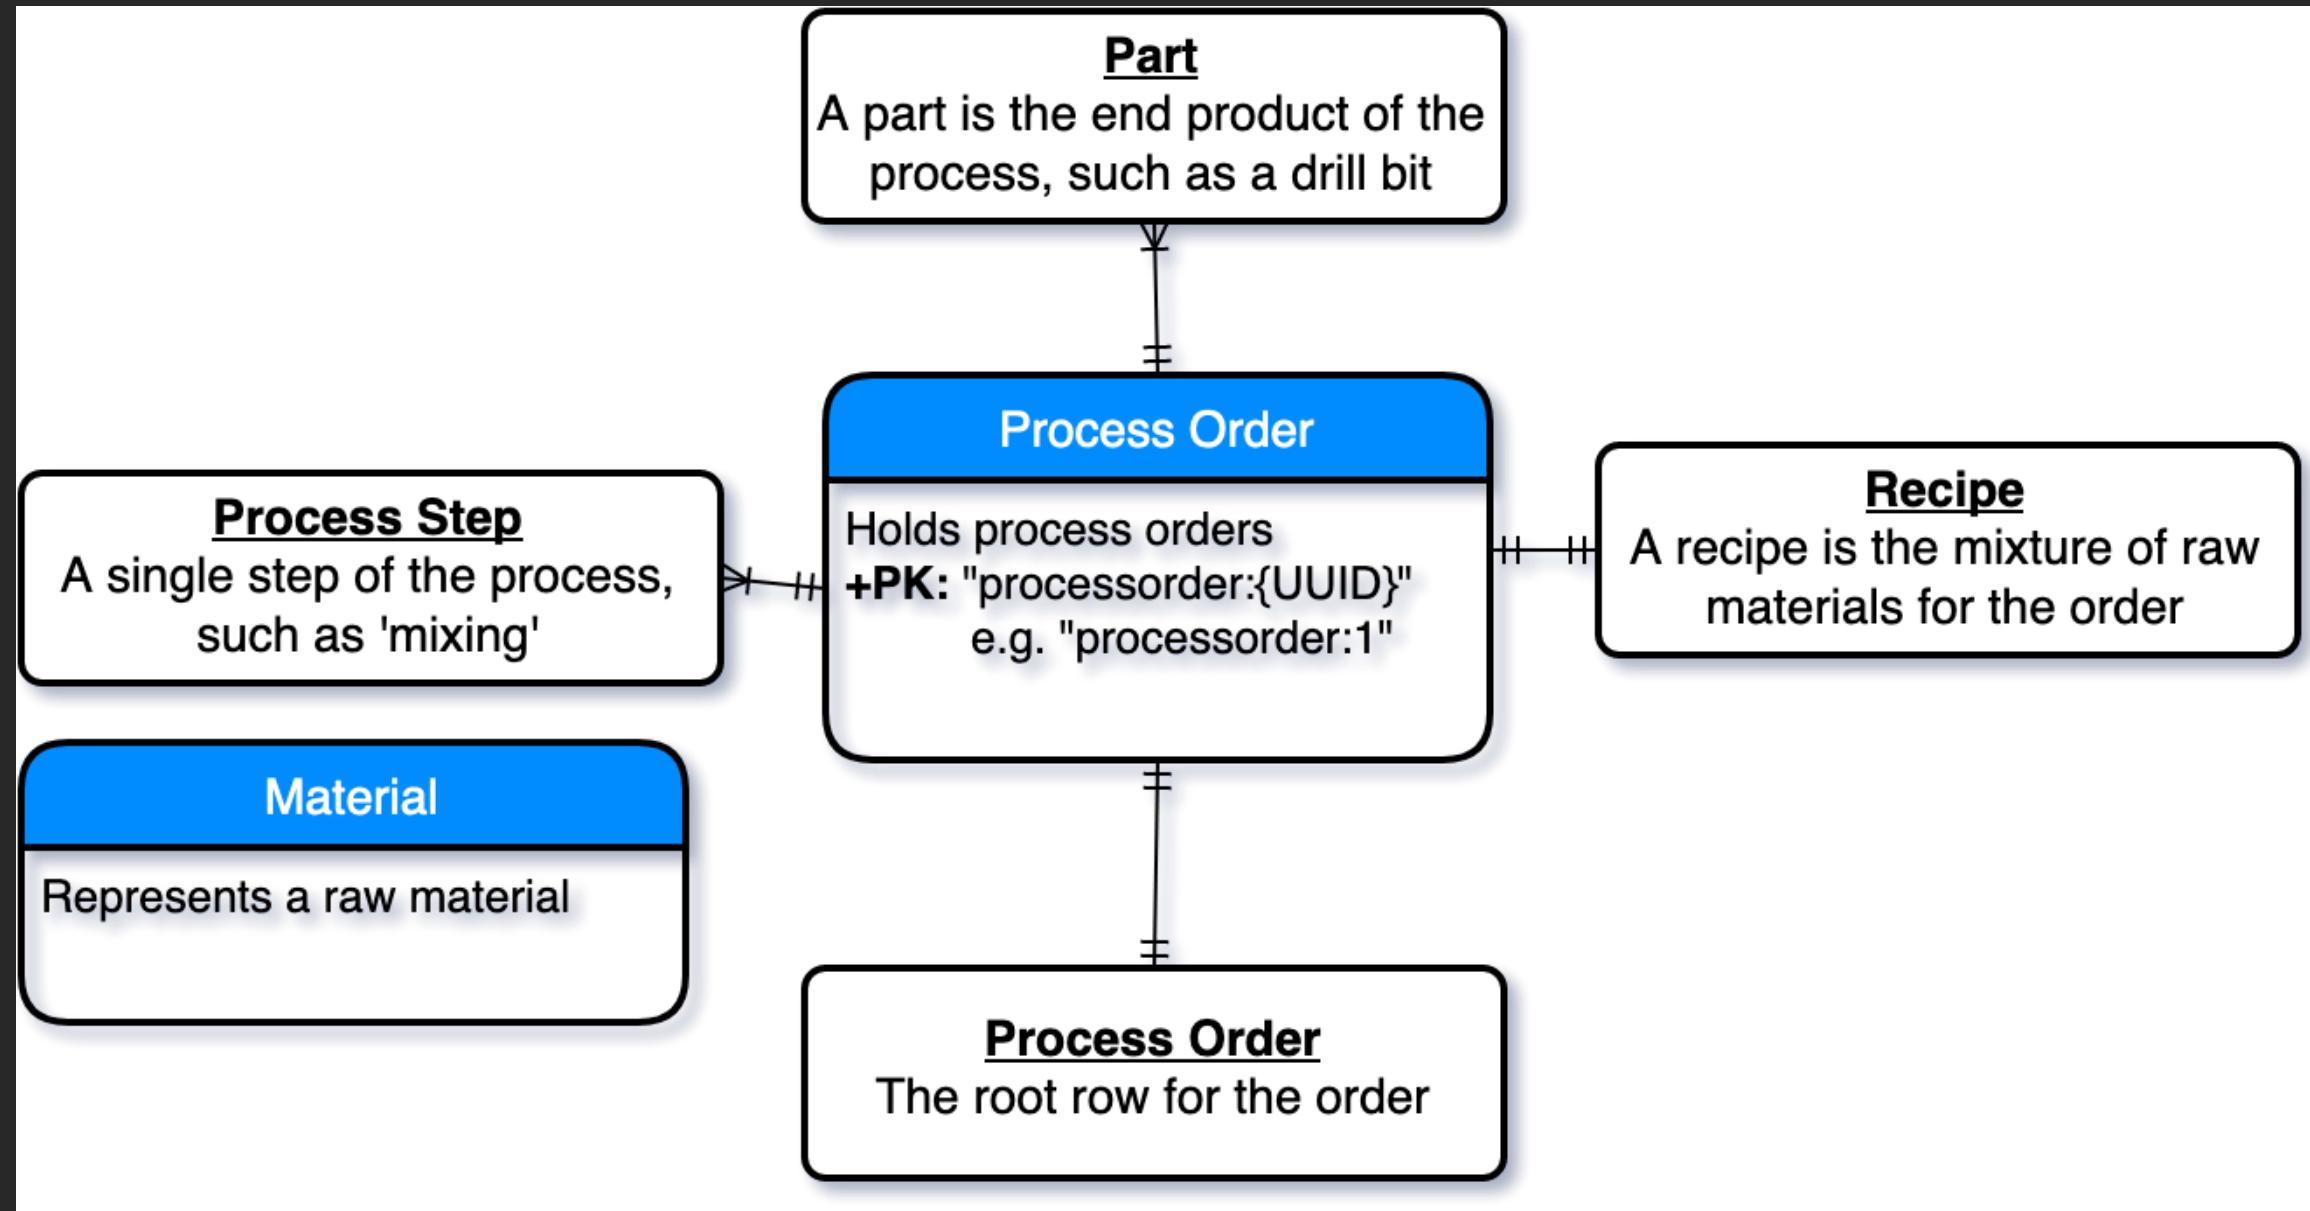
Step 4: Bake  
“Parts”

*“Light the fires!”*

Step 5: QC  
“Parts”

*Microscopy  
Measurement*

# Establish the entity relationships



# Queries

- ✓ Get all steps for a PROCESS ORDER
- ✓ Get the AVERAGE density of a MATERIAL
- ✓ Find the PROCESS ORDER for a PART

# Final schema

Primary Key		Attributes		
PK	SK			
processorder:1	part:1	porosity	mass	GSI_1_PK
		20%	122g	part:1
	processstep:materials:1	materialIds		
		[materialId:1, materialId:2]		
	processstep:mixing:2	mass	volume	density
		130g	26cm^3	5g/cm^3
	processstep:print:3	bedTemp	materialTemp	adhesionSurface
		120	400	tape
	recipes	processstep:materials		
		recipe:1		

Primary Key		Attributes			
PK	SK				
materialid:1	processorder:1	mass	volume	density	timestamp
		45g	9cm^3	5g/cm^3	1575423000
materialid:2	processorder:2	mass	volume	density	timestamp
		66g	13.2cm^3	5g/cm^3	1575477000
materialid:1	processorder:1	mass	volume	density	timestamp
		85g	17cm^3	5g/cm^3	1575423000
	processorder:2	mass	volume	density	timestamp
		64g	12.8cm^3	5g/cm^3	1575477000

# Adjacency lists

**Generic partition and sort key names**

PK	SK	Attributes			
	part:1	porosity	20%	mass	122g
processorder:1	processstep:materials:1	materialIds	[materialId:1, materialId:2]	volume	26cm^3
	processstep:mixing:2	mass	130g	density	5g/cm^3
	processstep:print:3	bedTemp	120	materialTemp	400
		adhesionSurface			
		tape			

**Often used with GSI overloading**

GSI_1_PK
part:1

**Soft associations for n:m relationship**

Primary Key		Attributes			
PK	SK	mass	volume	density	timestamp
materialid:1	processorder:1	45g	9cm^3	5g/cm^3	1575423000
	processorder:2	66g	13.2cm^3	5g/cm^3	1575477000
materialid:2	processorder:1	85g	17cm^3	5g/cm^3	1575423000
	processorder:2	64g	12.8cm^3	5g/cm^3	1575477000

# Queries

Primary Key		Attributes		
PK	SK			
processorder:1	part:1	porosity	mass	GSI_1_PK
		20%	122g	part:1
	processstep:materials:1	materialIds		
		[materialId:1, materialId:2]		
	processstep:mixing:2	mass	volume	density
		130g	26cm^3	5g/cm^3
	processstep:print:3	bedTemp	materialTemp	adhesionSurface
		120	400	tape
	recipes	processstep:materials		
		recipe:1		

Primary Key		Attributes		
PK				
materialid:1		Query #1: Get all steps for a PROCESS ORDER		
				timestamp 1575423000
materialid:2		Solution: Query with key_condition		
		PK = process_order:1 AND SK starts_with("processstep")		
				timestamp 1575477000
				timestamp 1575423000
				timestamp 1575477000

# Queries

Primary Key		Attributes	
PK	SK	GSI_1_PK	part:1
processor	Query #2: Get the AVERAGE density of a MATERIAL	density	5g/cm <sup>3</sup>
processor	Solution: Query with key_condition PK = materialId:1 and on the client compute the AVERAGE	adhesionSurface	tape
recipes	recipe:1		

Primary Key		Attributes			
PK	SK	mass	volume	density	timestamp
materialid:1	processorder:1	45g	9cm^3	5g/cm^3	1575423000
		66g	13.2cm^3	5g/cm^3	1575477000
	processorder:2	85g	17cm^3	5g/cm^3	1575423000
		64g	12.8cm^3	5g/cm^3	1575477000
materialid:2	processorder:1	mass	volume	density	timestamp
		85g	17cm^3	5g/cm^3	1575423000
	processorder:2	mass	volume	density	timestamp
		64g	12.8cm^3	5g/cm^3	1575477000

# Queries

Primary Key		Attributes			
PK	SK				
processorder:1	part:1	porosity	mass	GSI_1_PK	
		20%	122g	part:1	
	processstep:materials:1	materialIds			
		[materialId:1, materialId:2]			
	processstep:mixing:2	mass	volume	density	
		130g	26cm^3	5g/cm^3	
	processstep:print:3	bedTemp	materialTemp	adhesionSurface	
		120	400	tape	
	recipes	processstep:materials			
		recipe:1			

Primary Key		Attributes			
PK	SK				
materialid:1					timestamp
					1575423000
materialid:2					timestamp
					1575477000
Process Order		64g	12.8cm^3	5g/cm^3	timestamp
					1575423000
					timestamp
					1575477000

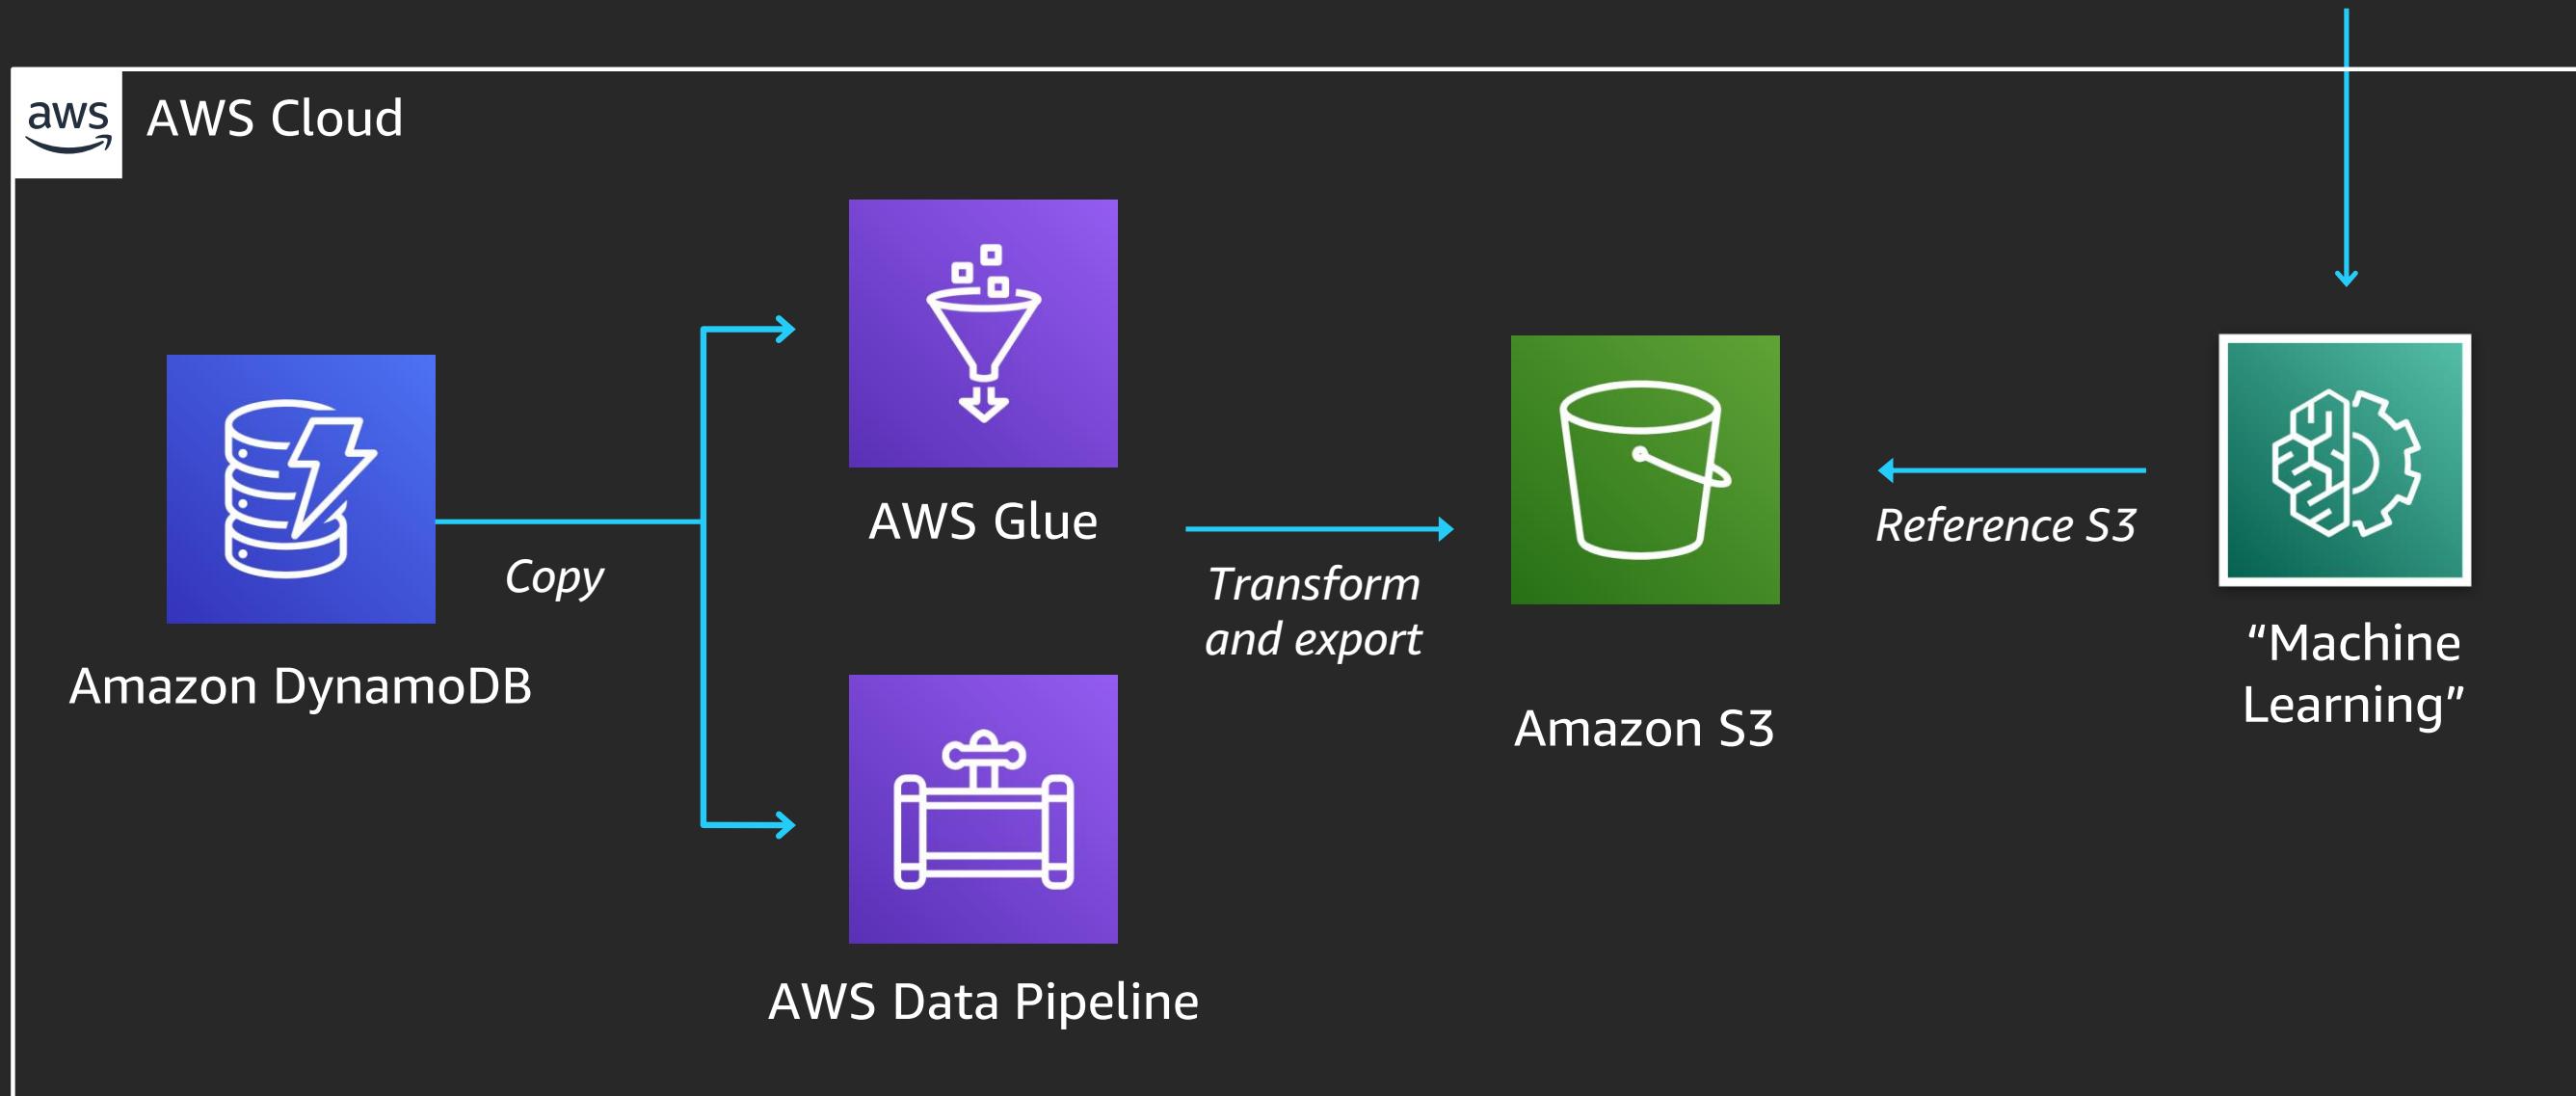
Query #3: Find the PROCESS ORDER for a PART

Solution: Query GSI 1 with key\_condition GSI\_1\_PK = part:1

# Data lake: Export to Amazon S3



Data  
scientists



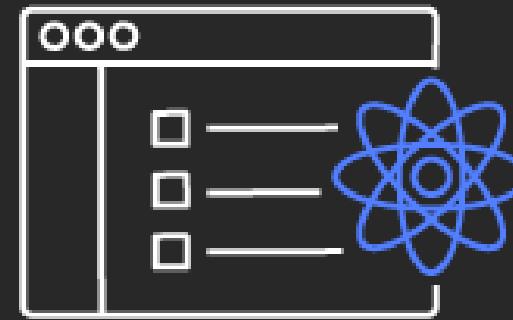
# Recap



✓ Adjacency lists

Good for:

- n:m relationships
- Relational data



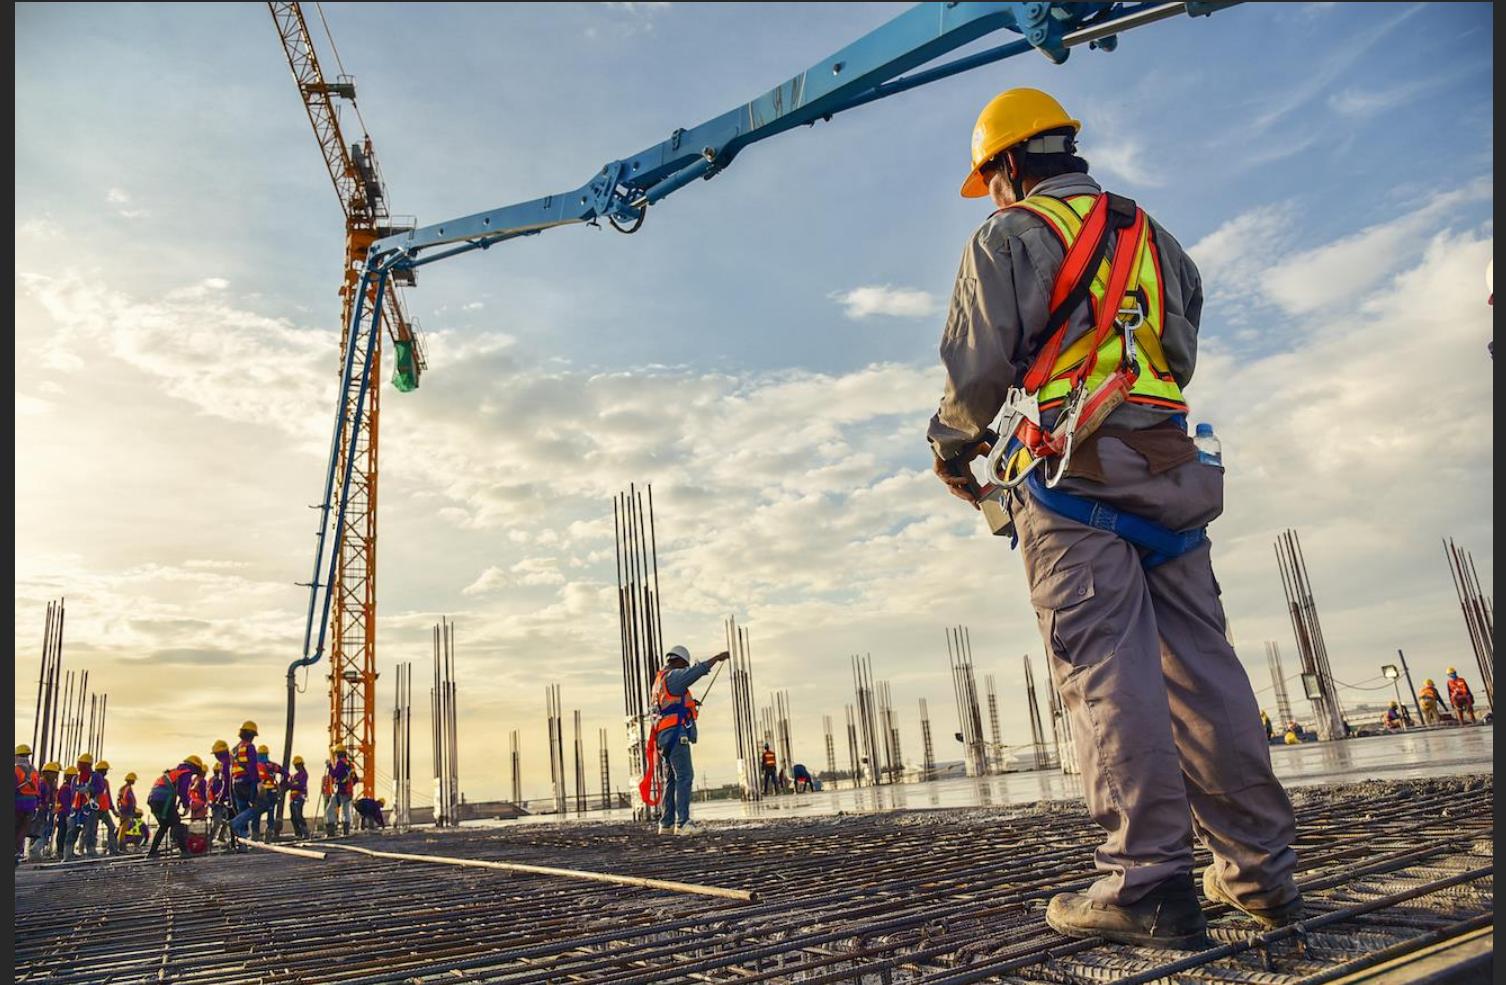
✓ Export data to S3 with AWS

Good for:

- Machine learning
- Data lakes

# Customer #2: Heavy machinery

*Use case: access control database, n:m consistency*



# Use case



User

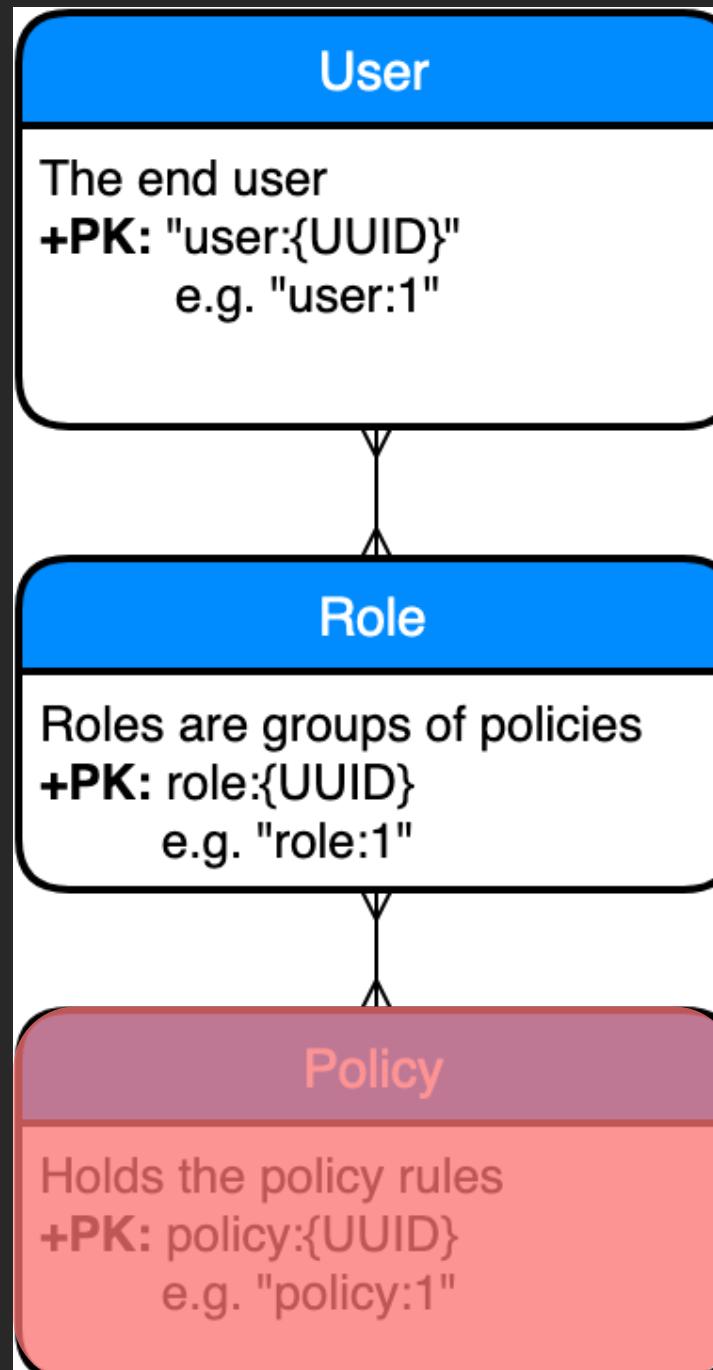


Role



Policy

# Entity relationships (original)



Get user → Get roles for user → Get policies for role

10ms per read, limited scalability

# All data is relational



IT monitoring

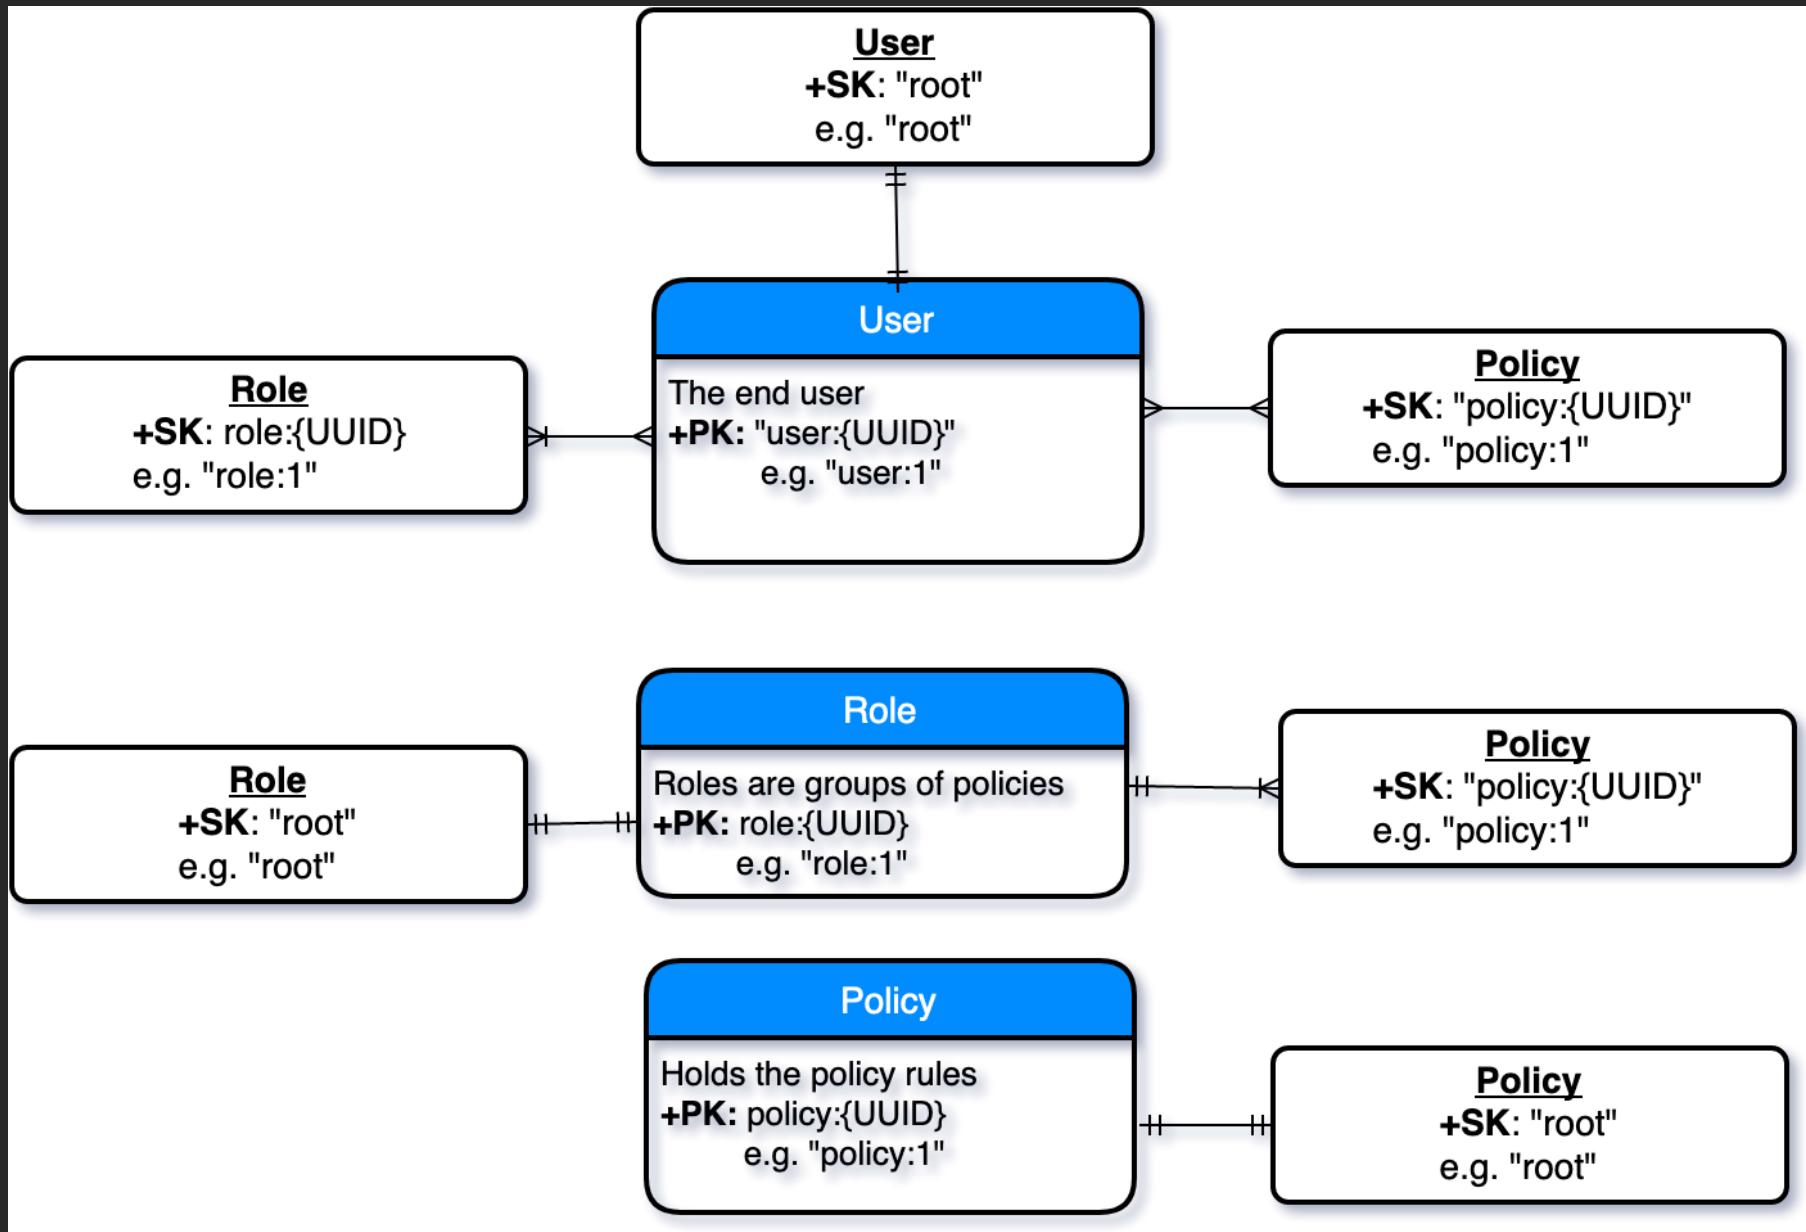


Social graph



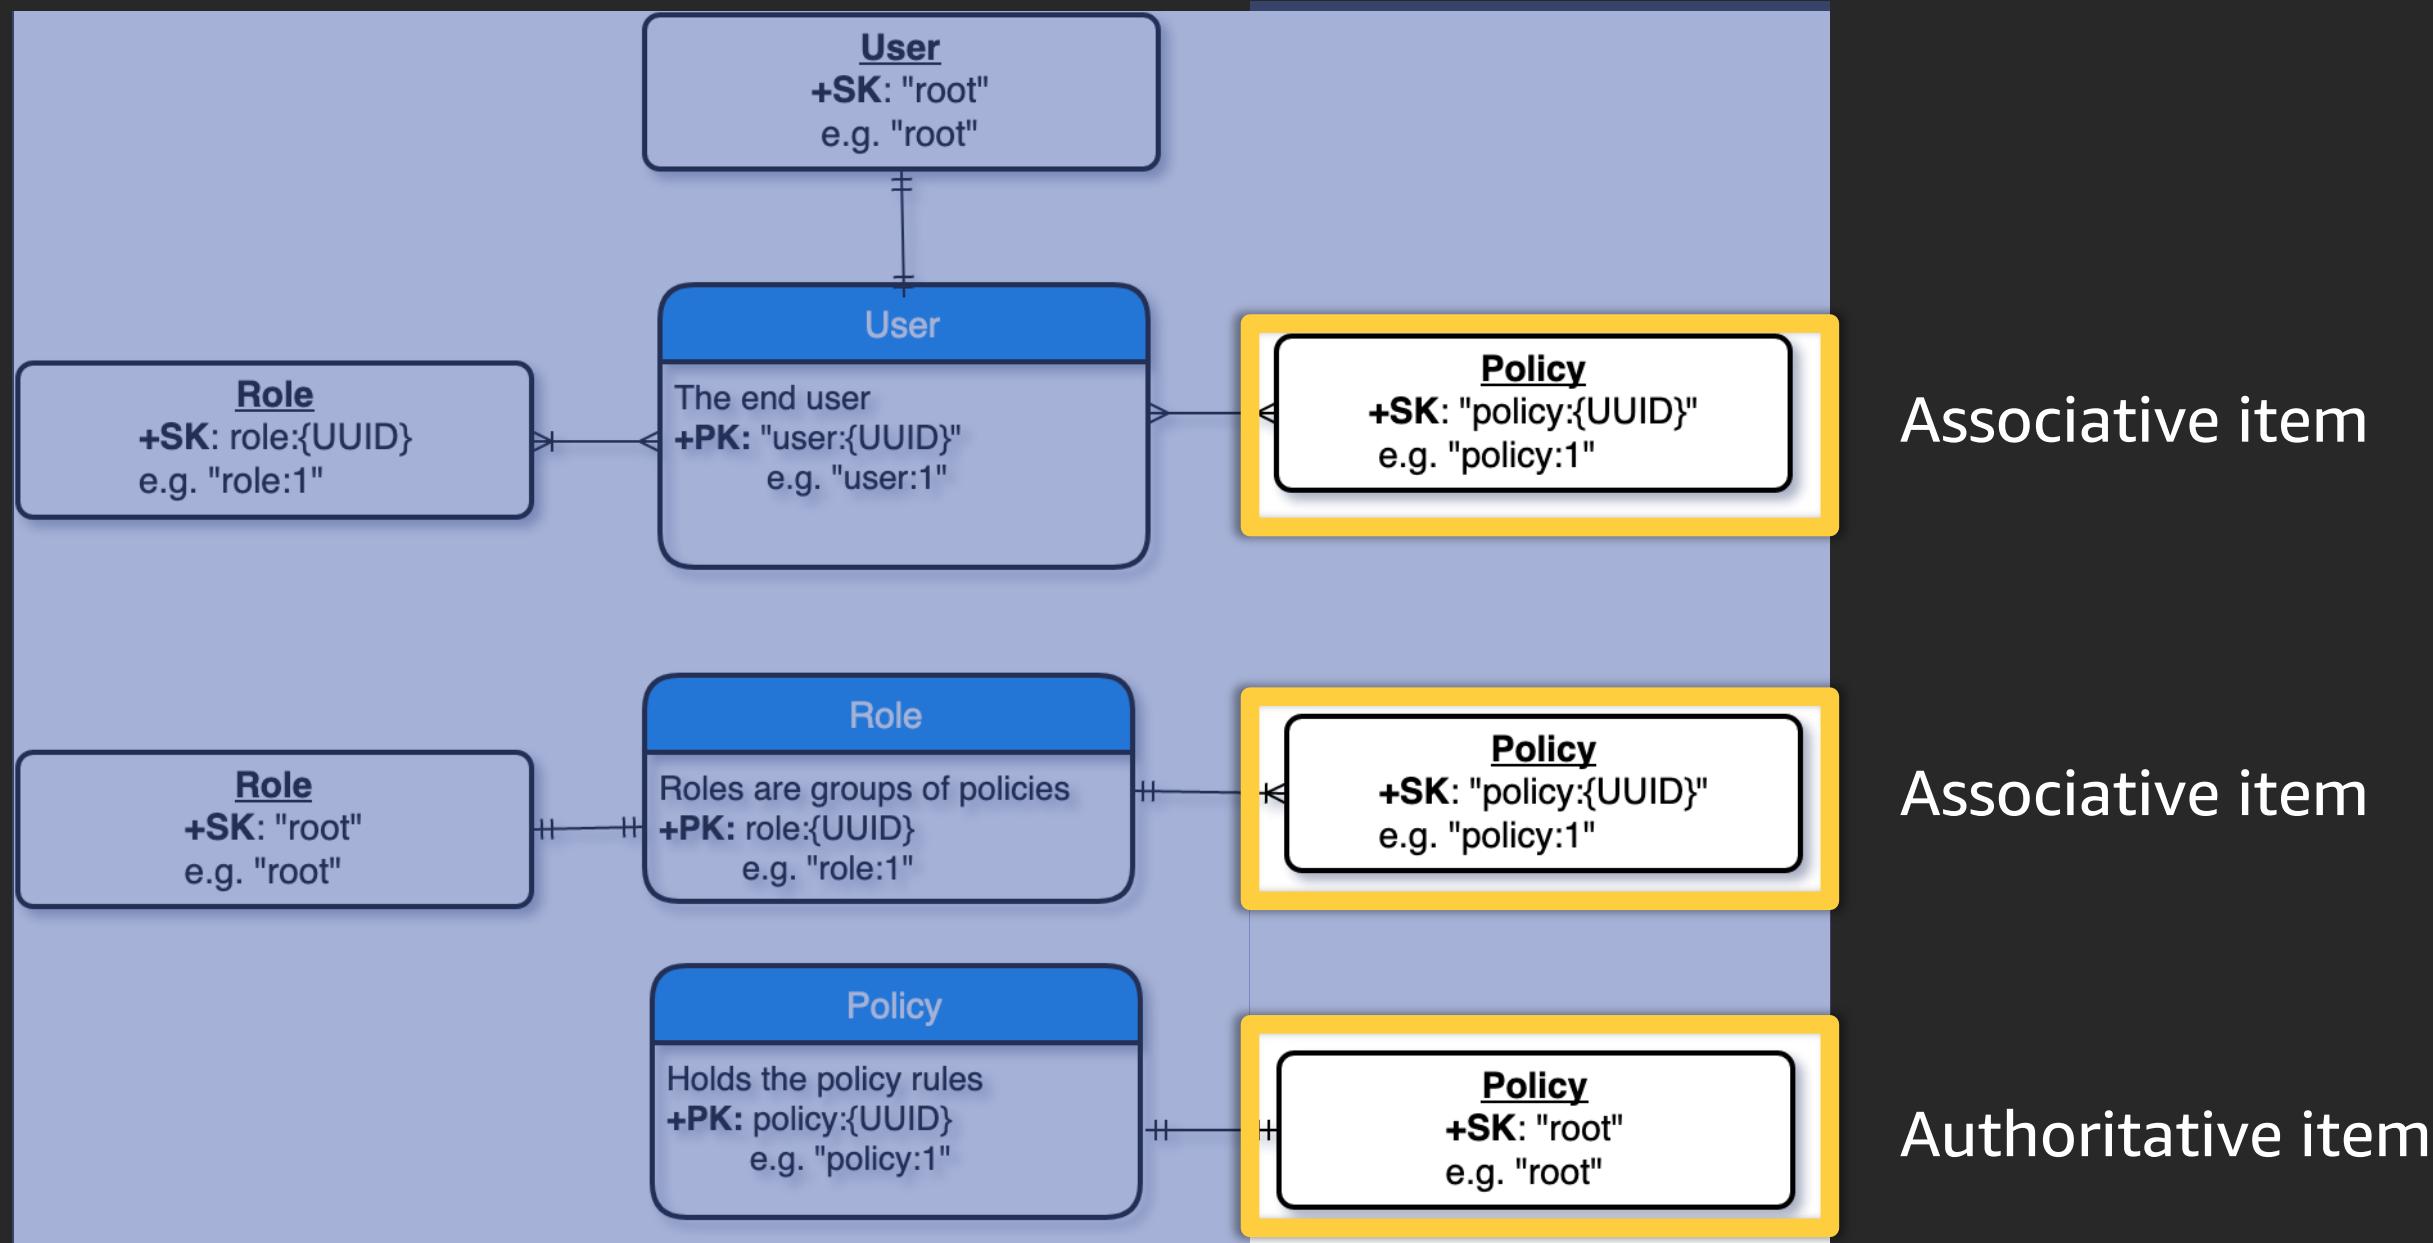
Documents

# Entity relationships (fixed)



Aggregation → scalability

# Entity relationships



Aggregation → scalability

# Queries

- ✓ Get all policies for a USER
- ✓ Find all POLICIES for a ROLE
- ✓ Find all USERS in an OrgId

# Final schema

Primary Key		Attributes	
PK	SK		
policy:1	root	policyName	policy
		Admin	{ ... }
		policyName	policy
			RejectWrites
		policyName	policy
			RevokeAssumeRole

Primary Key		Attributes	
PK	SK		
role:1	root	roleName	
		Default role	
		policyName	GSI_1_PK
			role:policy:1
		policyName	GSI_1_PK
			role:policy:2
policy:1	policy:1	policyName	GSI_1_PK
			role:policy:3
		policyName	
		RevokeAssumeRole	

Primary Key		Attributes		
PK	SK			
user:1	root	userName	orgId	GSI_1_PK
		John Doe	org:1	org:1
		roleName	GSI_1_PK	
			role:1	
		Default role		
policy:1	policy:1	policyName	policy	GSI_1_PK
			{ ... }	user:policy:1
		Admin		
		policyName	policy	GSI_1_PK
			{ ... }	user:policy:2
policy:2	policy:2	RejectWrites		
		policyName	policy	GSI_1_PK
			{ ... }	user:policy:3
		RevokeAssumeRole		

# Queries

Query #1: Get all policies for a USER

Solution: Query with key\_condition:

PK = user:1 AND SK starts\_with("policy")

Primary Key		Attributes		
PK	SK	userName	orgId	GSI_1_PK
user:1	root	John Doe	org:1	org:1
		roleName	GSI_1_PK	
	role:1	Default role	role:1	
		policyName	policy	GSI_1_PK
user:1	policy:1	Admin	{ ... }	user:policy:1
		policyName	policy	GSI_1_PK
	policy:2	RejectWrites	{ ... }	user:policy:2
		policyName	policy	GSI_1_PK
	policy:3	RevokeAssumeRole	{ ... }	user:policy:3

# Queries

Primary Key		Attributes	
PK	SK	roleName	
	root	Default role	
role:1	policy:1	policyName	GSI_1_PK
		Admin	role:policy:1
	policy:2	policyName	GSI_1_PK
		RejectWrites	role:policy:2
	policy:3	policyName	GSI_1_PK
		RevokeAssumeRole	role:policy:3

Query #2: Find all POLICIES for a ROLE

Solution: Query with key\_condition:

PK = role:1 AND SK starts\_with("policy")

# Queries

Query #3: Find all USERS in an OrgId

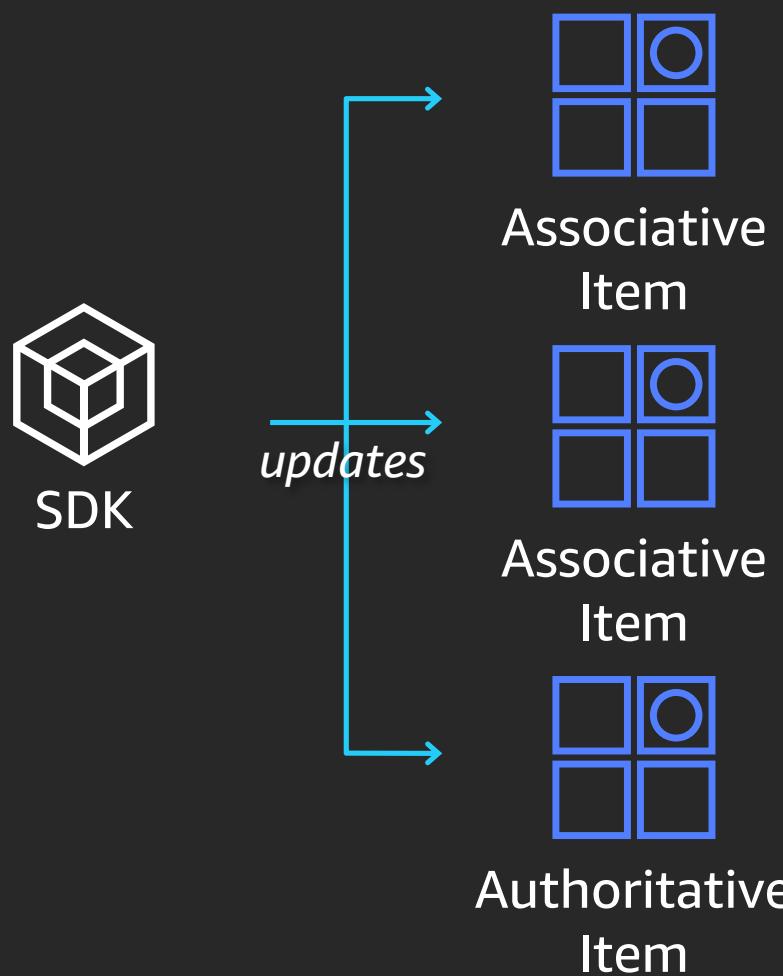
Solution: Query GSI 1 with key\_condition: PK = org:1

Note: The full row is projected into the GSI

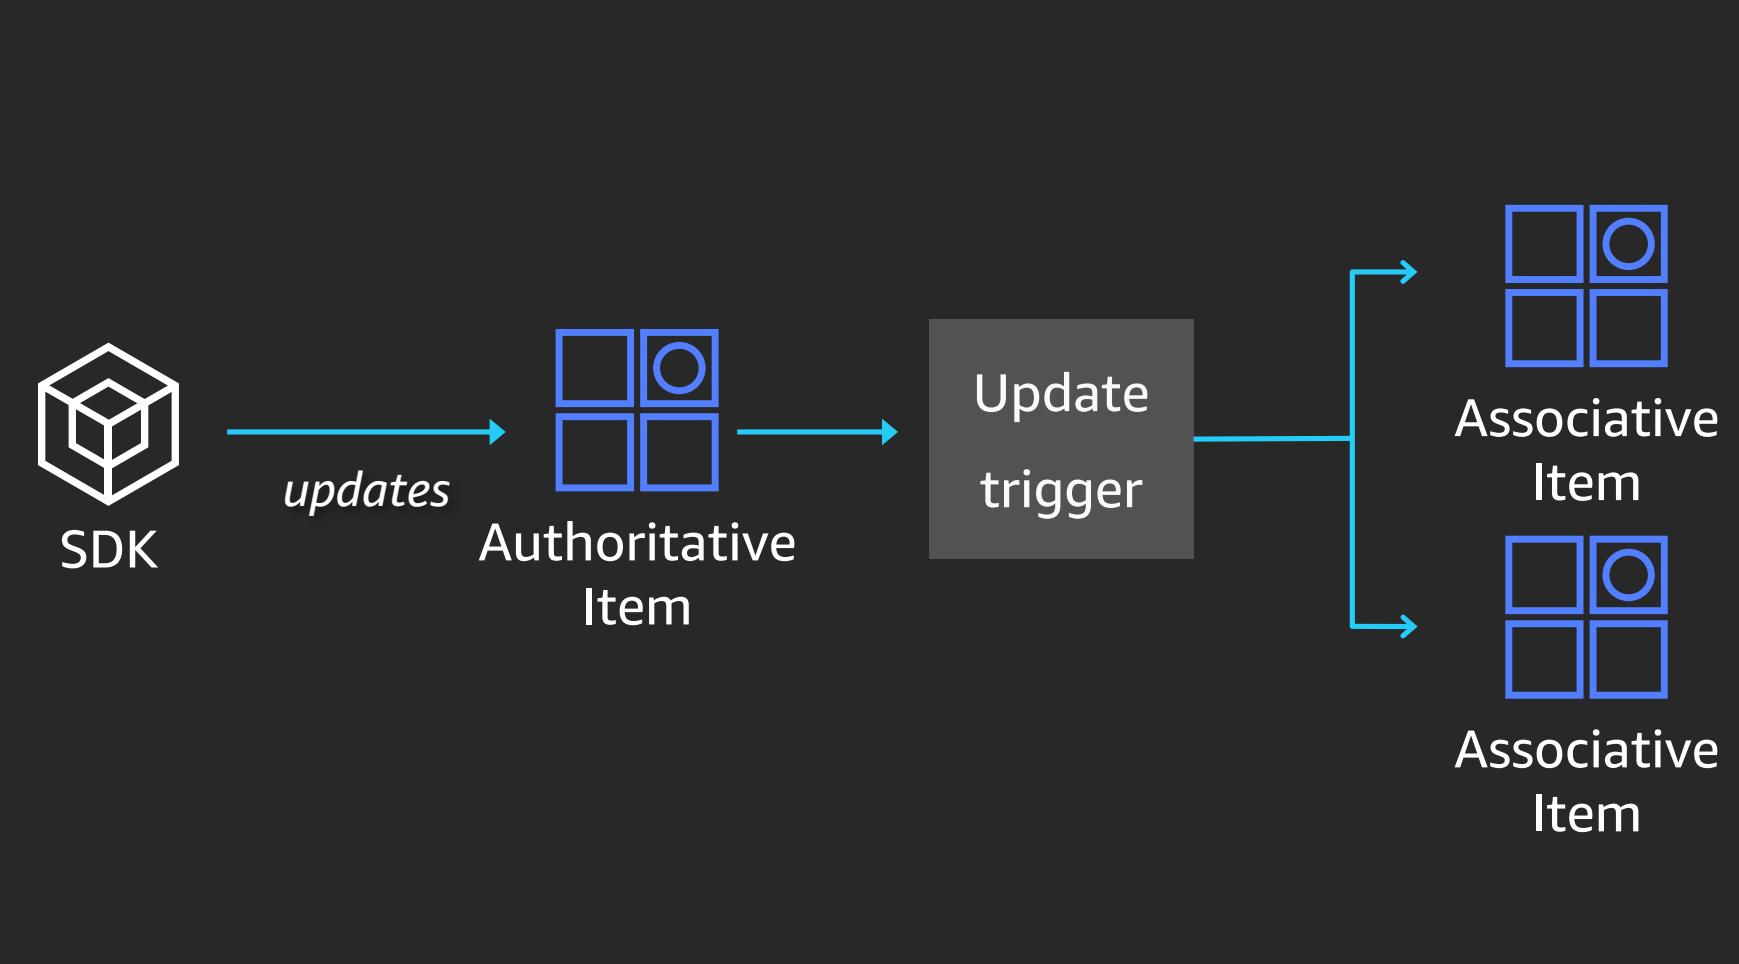
Primary Key		Attributes		
PK	SK			
	root	userName	orgId	GSI_1_PK
		John Doe	org:1	org:1
user:1	role:1	roleName	GSI_1_PK	
		Default role	role:1	
	policy:1	policyName	policy	GSI_1_PK
		Admin	{ ... }	user:policy:1
	policy:2	policyName	policy	GSI_1_PK
		RejectWrites	{ ... }	user:policy:2
	policy:3	policyName	policy	GSI_1_PK
		RevokeAssumeRole	{ ... }	user:policy:3

# Update strategies for references

## Immediate writes

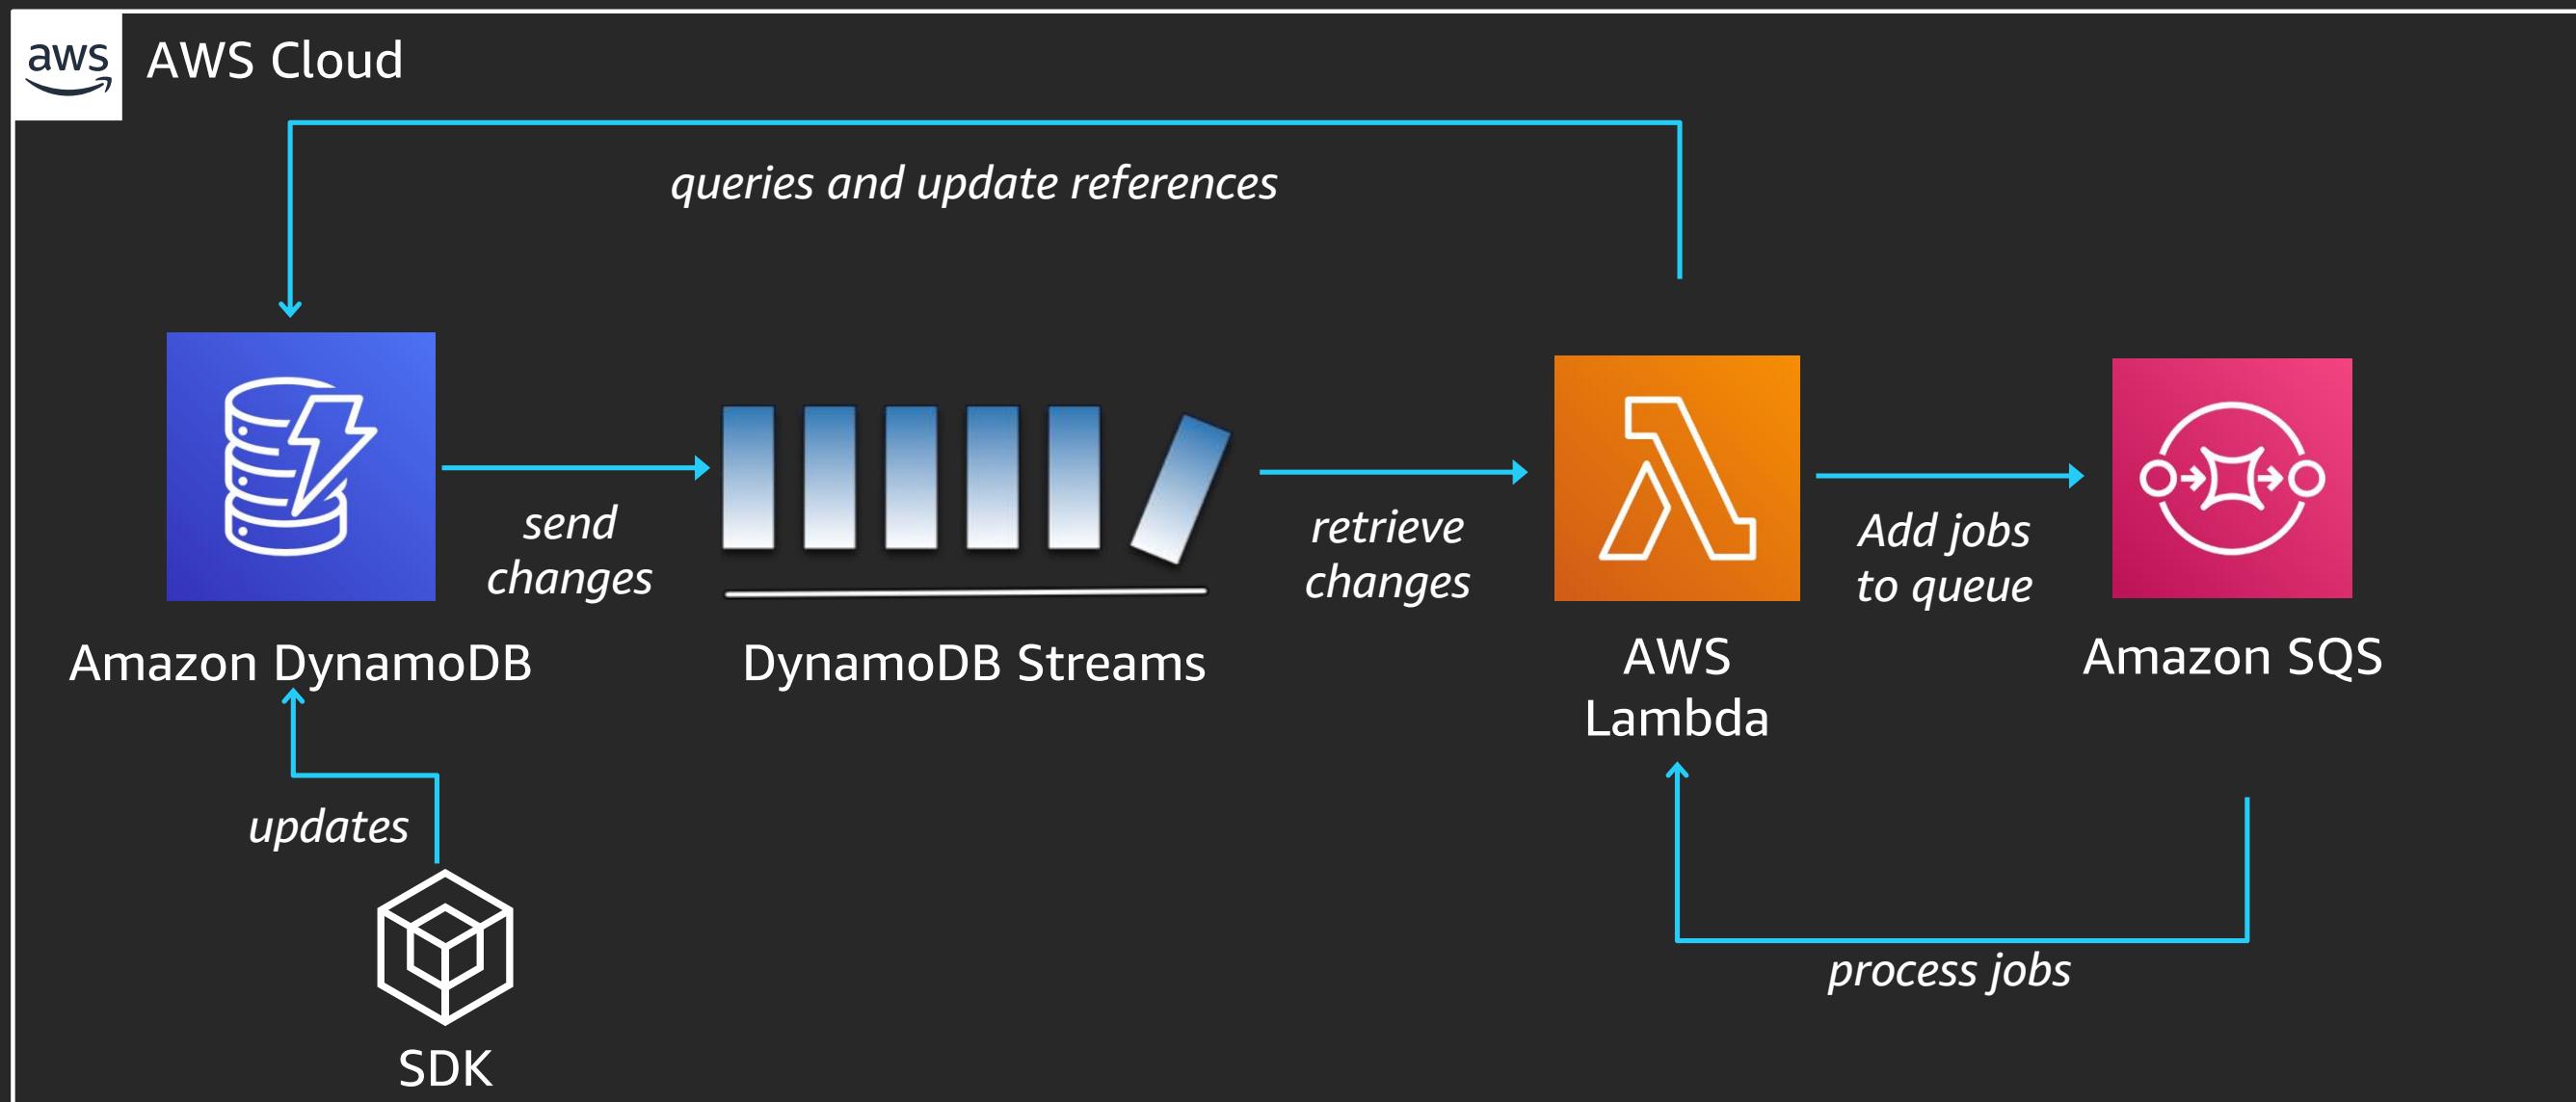


## Triggers with DynamoDB streams



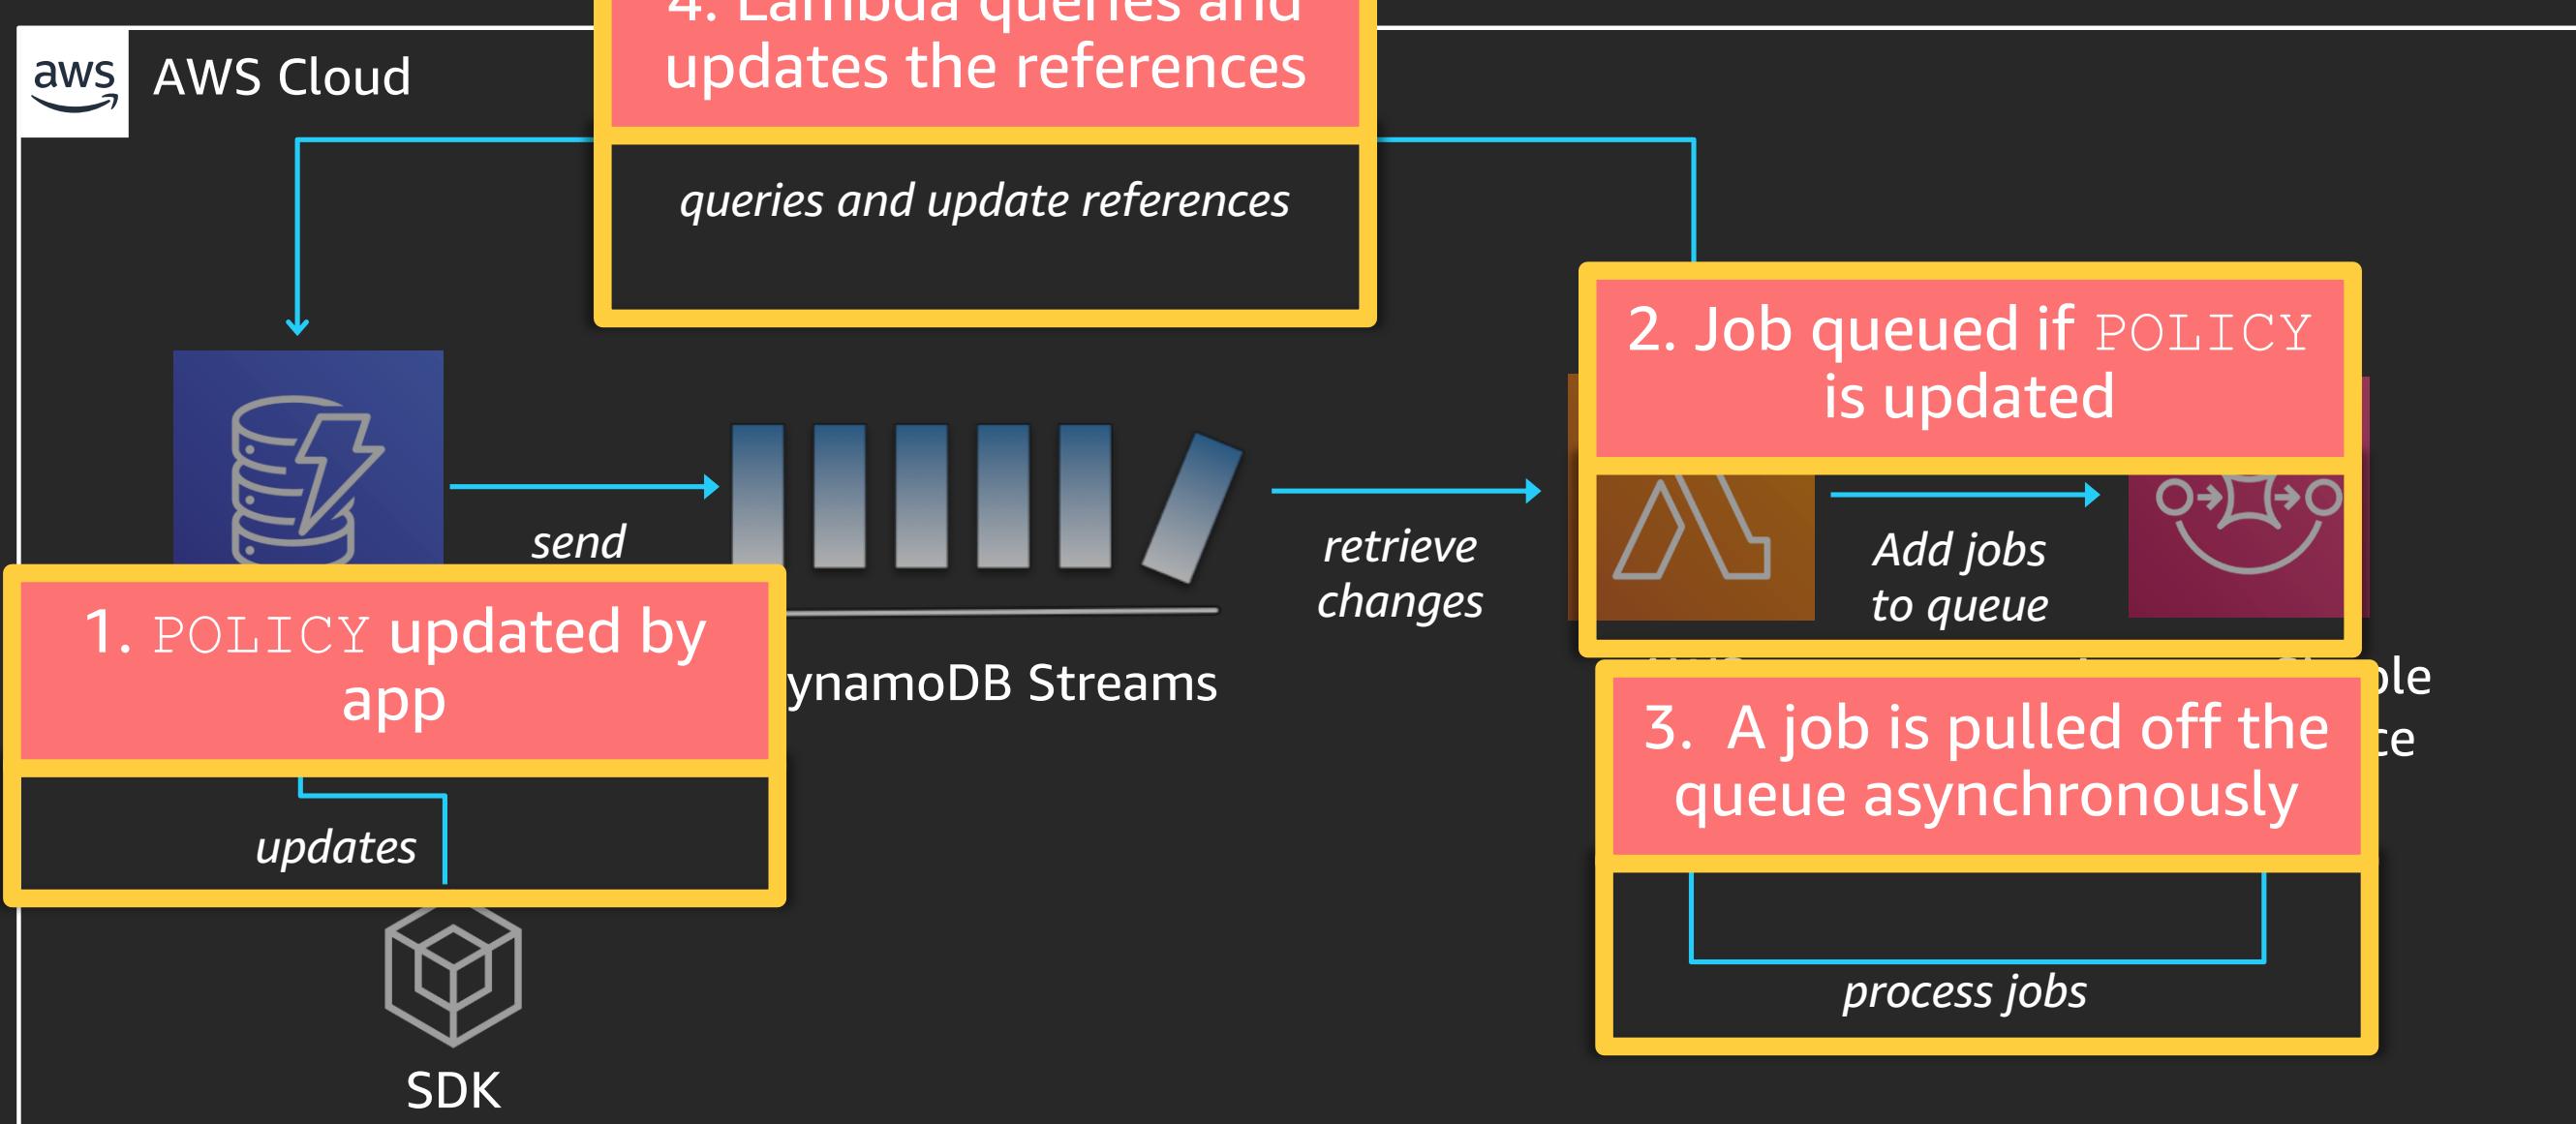
# Update triggers

1. Update a POLICY in DynamoDB
2. Queue a job to update the references
3. Process the job
4. Update the references



# Update triggers

1. Update a POLICY in DynamoDB
2. Queue a job to update the references
3. Process the job
4. Update the references



# Recap



✓ Aggregate for performance

Good for:

- Relational data



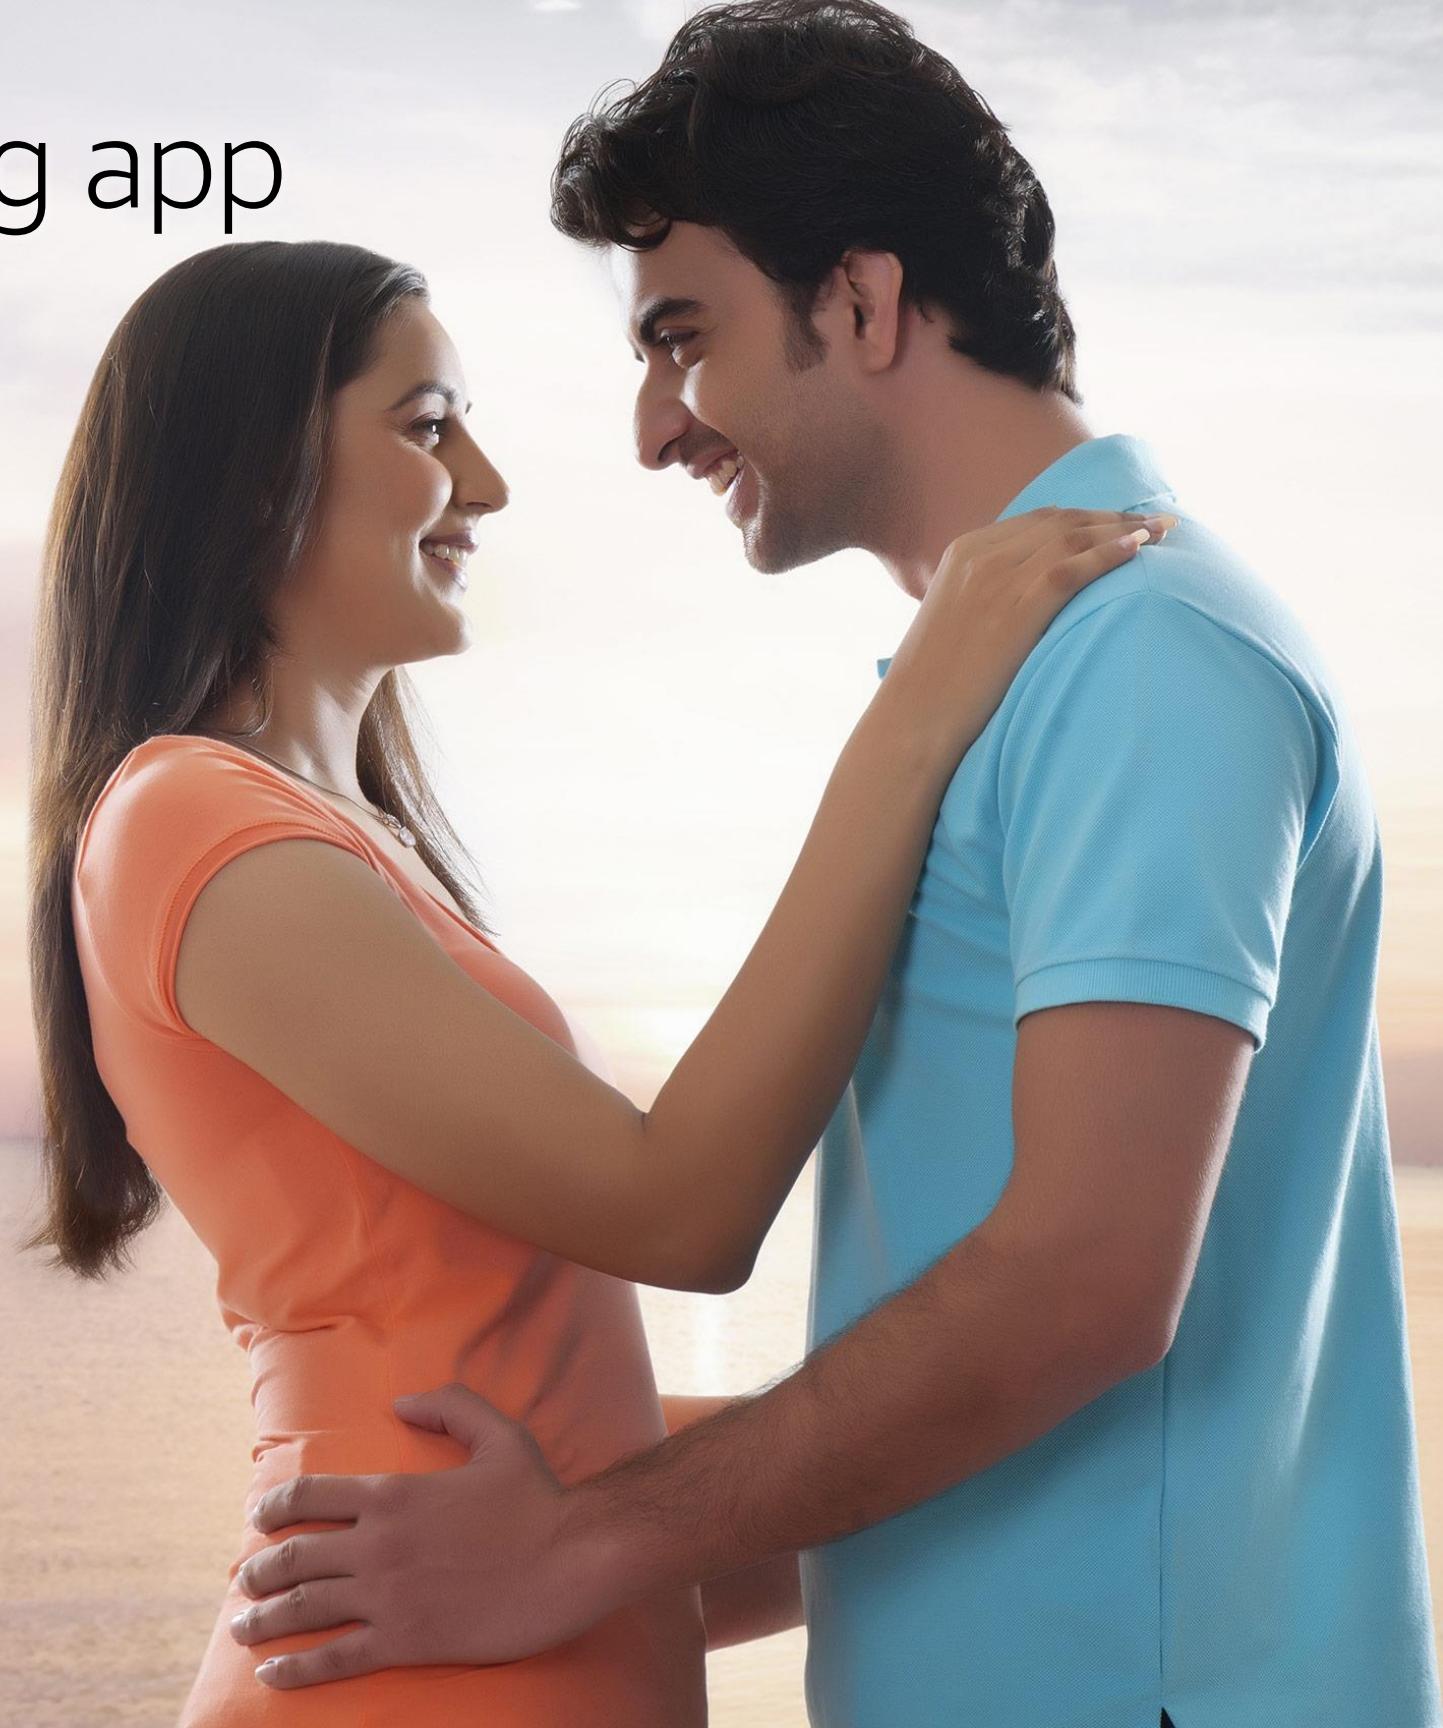
✓ Trigger based updates

Good for:

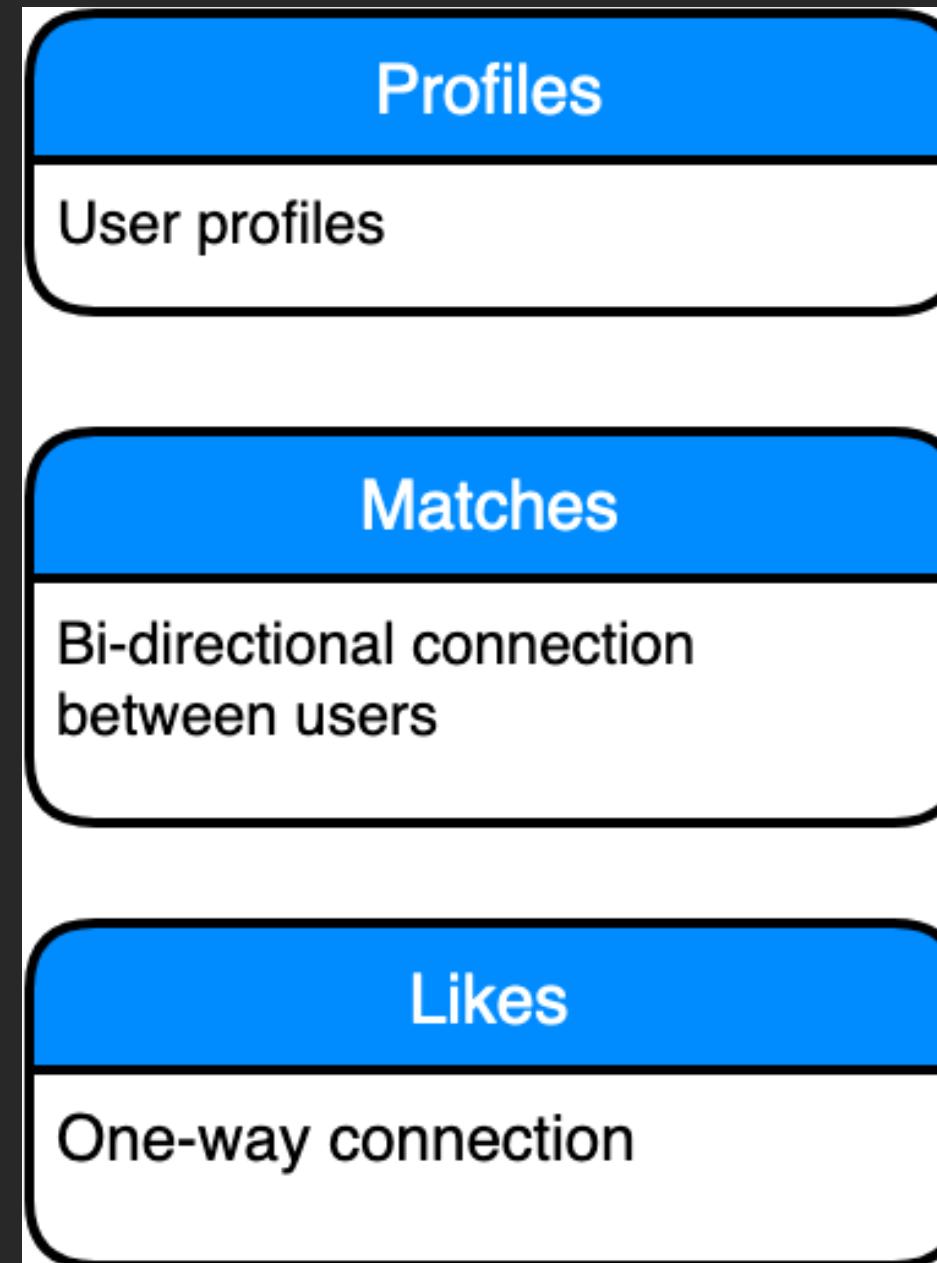
- Horizontal scaling
- Relational data

Customer #3: Dating app

*Use case: job scheduling*



# Entity relationships



# Queries

- ✓ Get MATCHES for a user
- ✓ For a user, find who LIKED them
- ✓ Find MATCHES that are stale

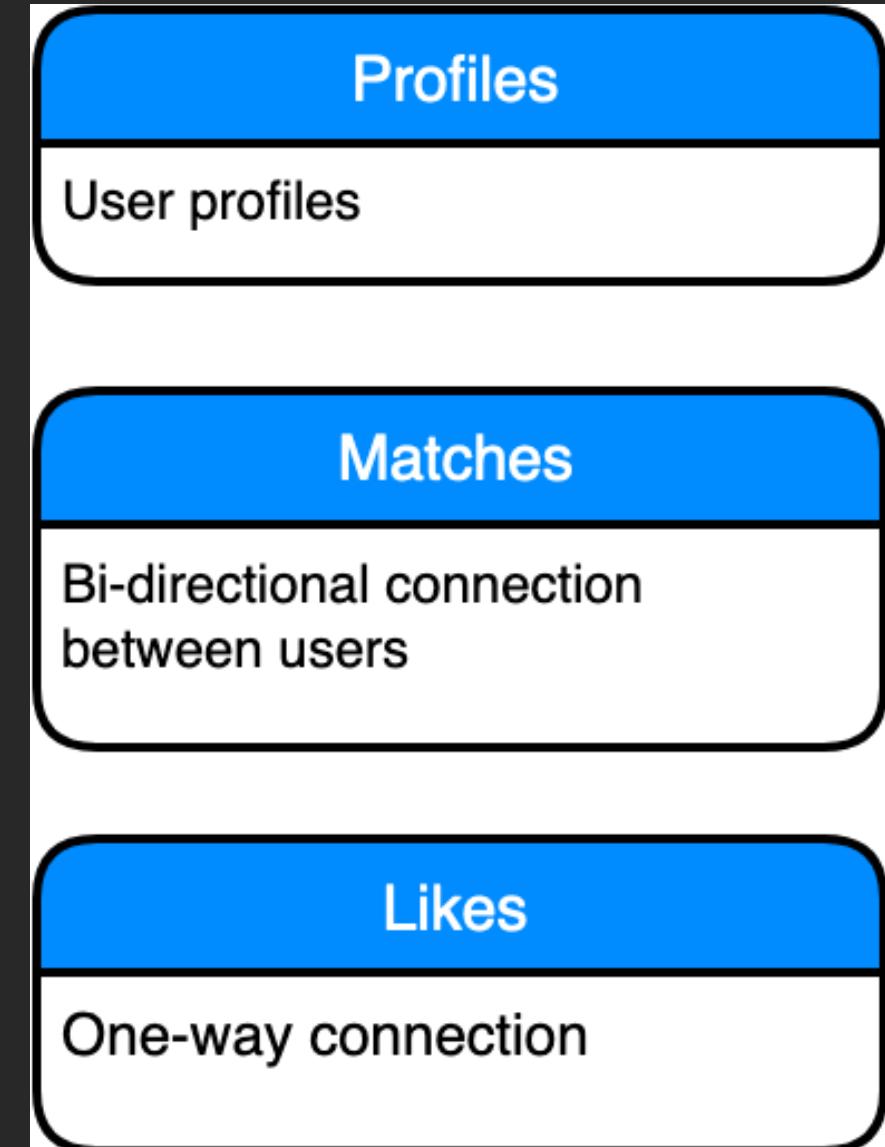
# Table schema



# All together, now

Keep in mind:

- Establish the primary keys first, and then add the GSIs
- Not all entities will be used to establish the queries
- The solution will use one or two queries
- A stale match is one that has not been seen by either party in the last 30 days



- ✓ Get MATCHES for a user
- ✓ For a user, find who LIKED them
- ✓ Find MATCHES that are stale

# Final table schema

Primary Key		Attributes					
PK	SK						
match:1	root	user1	user2	GSI_1_PK	GSI_2_PK	GSI_3_PK	GSI_3_SK
		user:1234	user:4567	user:1234	user:4567	Match#(0-N)	2019-12-31T03:44
match:2	root	user1	user2	GSI_1_PK	GSI_2_PK	GSI_3_PK	GSI_3_SK
		user:9833	user:4590	user:9833	user:4590	Match#(0-N)	2019-12-31T03:44

Primary Key		Attributes			
PK	SK				
user:1234	user:4567	Status	deleted	GSI_1_PK	GSI_1_SK
		active	FALSE	user:like:4567	1577829985
user:4567	user:1234	Status	deleted	GSI_1_PK	GSI_1_SK
		active	FALSE	user:like:1234	1577833583

# Table schema

Query #1: Get MATCHES for a user  
Solution: : Query GSI 1 with key\_condition  
GSI\_PK\_1 = user:1234, Query GSI 2 with  
key\_condition GSI\_PK\_2 = user:1234

Primary Key		Attributes					
PK	SK	#1			#3		
match:1	root	user1	user2	GSI_1_PK	GSI_2_PK	GSI_3_PK	GSI_3_SK
		user:1234	user:4567	user:1234	user:4567	Match#(0-N)	2019-12-31T03:44
match:2	root	user1	user2	GSI_1_PK	GSI_2_PK	GSI_3_PK	GSI_3_SK
		user:9833	user:4590	user:9833	user:4590	Match#(0-N)	2019-12-31T03:44

Primary Key		Attributes				
PK	SK	#2				
user:1234	user:4567	Status	deleted	GSI_1_PK	GSI_1_SK	
		active	FALSE	user:like:4567	577829985	
user:4567	user:1234	Status	deleted	GSI_1_PK	GSI_1_SK	
		active	FALSE	user:like:1234	1577833583	

Query #2: For a user, find who LIKED them  
Solution: : Query GSI 1 with key\_condition  
GSI\_PK\_1 = user::like:1234

Query #3: Find MATCHES that are stale  
Solution: : Query every Match shard of GSI 3  
with key\_condition GSI\_PK\_1 =  
Shard# (SHARDNUMBER) AND GSI\_1\_SK  
< 2020-01-15

# GSI overloading

Reuse generic index keys for different access patterns in the same table

Primary Key		Attributes			
PK	SK	JobName	InputType	GSI_1_PK	GSI_1_SK
job:1	root	"Transcribe audio"	mp3	Shard#(0-N)	2019-12-31
		"Verbatim"	wav	Shard#(0-N)	2019-12-29
job:2	root	JobName	InputType	GSI_1_PK	GSI_1_SK
		"Verbatim"	wav	Shard#(0-N)	2019-12-29

# Recap



✓ Migrate to DynamoDB

Good for:

- Operational workloads



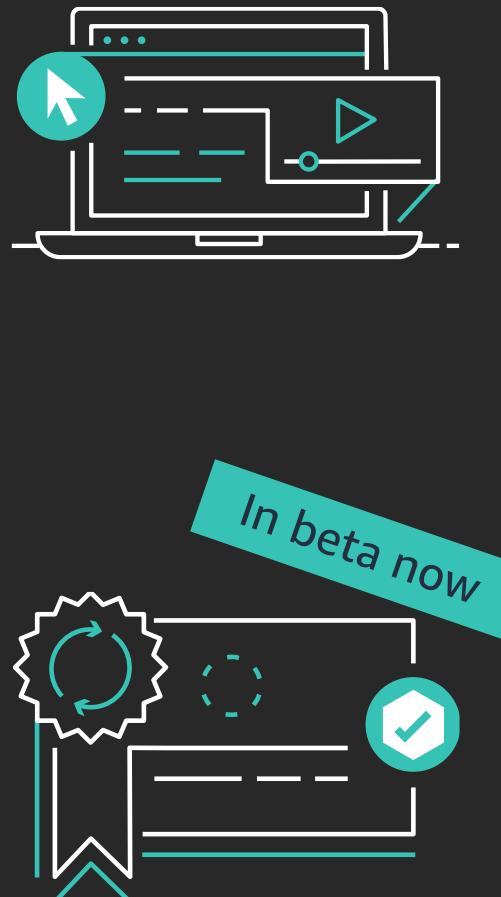
✓ Scalable search with GSI write sharding

Good for:

- Scheduled jobs
- Custom expiry rules

# Learn databases with AWS Training and Certification

Resources created by the experts at AWS to help you build and validate database skills



25+ free digital training courses cover topics and services related to databases, including:

- Amazon Aurora
- Amazon Neptune
- Amazon DocumentDB
- Amazon DynamoDB
- Amazon ElastiCache
- Amazon Redshift
- Amazon RDS

Validate expertise with the new **AWS Certified Database - Specialty** beta exam

[Visit aws.training](https://aws.training)

# Thank you!

**Sean Shriver & Ankur Kasliwal**

@sean\_shriver, @ankurrk, and  
@DynamoDB



Please complete the session  
survey in the mobile app.